

BCGX_churn

July 14, 2025

1 BCGX PowerCo Churn Analysis and Modeling

This project was completed as a part of BCG Data Science Job Simulation. Task for this project is to assist client PowerCo in churn analysis and modeling and analyze whether Price Sensitivity is a major contributing factor for Customer Churn. Price Sensitivity is defined as Change in Demand/Change in Price for this project. Throughout the project we followed BCGX guidance and completed steps as a Data Scientist would during the job at BCG.

Besides that, we have performed our own additional EDA and Churn Modeling, with our best model better than BCGX Model Answer by several times, five times better in F1 Score and 10 times better in recall performance as BCGX used a simple Random Forest while we used Gradient Boosting algorithm with hyper parameter tuning, Dense Neural Network, and a BiDirectional LSTM-CNN models. Among them XGBoost model performed best and this corresponds to the industry norm where gradient boosting and other boosting algorithms worked better in structured tabular dataset, provided that we have limited data.

This project will be divided into following steps: 1. EDA and Feature Engineering 2. Model Training, Hyperparameter Tuning, and Evaluation 3. Conclusion

1.1 EDA and Feature Engineering

-
1. Import packages
 2. Load data
 3. Feature engineering alongside indepth Analysis
-

1.2 1. Import packages

```
[1]: #required packages

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.metrics import roc_curve, roc_auc_score, classification_report, \
    ↪confusion_matrix, f1_score
from sklearn.model_selection import GridSearchCV
```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
import tensorflow as tf
from tensorflow.keras.layers import Dense
from tensorflow.keras import Sequential
from tensorflow.keras.initializers import HeNormal
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Activation
from tensorflow.keras.layers import Masking, Bidirectional, LSTM, Conv1D, GlobalMaxPooling1D, Dense, Dropout, Input
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping

import warnings
warnings.filterwarnings('ignore')

```

1.3 2. Load data

We have been given two datasets for this task. One contain Clients data and other contains Monthly Price data for clients over an year. Through out the EDA we will cover all the necessary columns for both datasets.

```
[ ]: EDA and Feature Engineering
```

```

[2]: df = pd.read_csv('./clean_data_after_eda.csv')
df["date_activ"] = pd.to_datetime(df["date_activ"], format='%Y-%m-%d')
df["date_end"] = pd.to_datetime(df["date_end"], format='%Y-%m-%d')
df["date_modif_prod"] = pd.to_datetime(df["date_modif_prod"], format='%Y-%m-%d')
df["date_renewal"] = pd.to_datetime(df["date_renewal"], format='%Y-%m-%d')

```

```
[3]: df.head()
```

```

[3]:
      id                                     channel_sales \
0  24011ae4ebbe3035111d65fa7c15bc57  foosdfpfkusacimwkcsosbicdxkicaua
1  d29c2c54acc38ff3c0614d0a653813dd                                MISSING
2  764c75f661154dac3a6c254cd082ea7d  foosdfpfkusacimwkcsosbicdxkicaua
3  bba03439a292a1e166f80264c16191cb  lmkebamcaaclubfxadlmueccxoimlema
4  149d57cf92fc41cf94415803a877cb4b                                MISSING

      cons_12m  cons_gas_12m  cons_last_month  date_activ  date_end  \
0           0         54946                0  2013-06-15  2016-06-15
1         4660              0                0  2009-08-21  2016-08-30
2          544              0                0  2010-04-16  2016-04-16
3         1584              0                0  2010-03-30  2016-03-30

```

```
4      4425      0      526 2010-01-13 2016-03-07
```

```

date_modif_prod date_renewal forecast_cons_12m ... \
0      2015-11-01  2015-06-23      0.00 ...
1      2009-08-21  2015-08-31     189.95 ...
2      2010-04-16  2015-04-17      47.96 ...
3      2010-03-30  2015-03-31     240.04 ...
4      2010-01-13  2015-03-09     445.75 ...

```

```

var_6m_price_off_peak_var var_6m_price_peak_var \
0      0.000131      4.100838e-05
1      0.000003      1.217891e-03
2      0.000004      9.450150e-08
3      0.000003      0.000000e+00
4      0.000011      2.896760e-06

```

```

var_6m_price_mid_peak_var var_6m_price_off_peak_fix \
0      9.084737e-04      2.086294
1      0.000000e+00      0.009482
2      0.000000e+00      0.000000
3      0.000000e+00      0.000000
4      4.860000e-10      0.000000

```

```

var_6m_price_peak_fix var_6m_price_mid_peak_fix var_6m_price_off_peak \
0      99.530517      44.235794      2.086425
1      0.000000      0.000000      0.009485
2      0.000000      0.000000      0.000004
3      0.000000      0.000000      0.000003
4      0.000000      0.000000      0.000011

```

```

var_6m_price_peak var_6m_price_mid_peak churn
0      9.953056e+01      4.423670e+01      1
1      1.217891e-03      0.000000e+00      0
2      9.450150e-08      0.000000e+00      0
3      0.000000e+00      0.000000e+00      0
4      2.896760e-06      4.860000e-10      0

```


```
[5 rows x 44 columns]
```

```
[4]: #We will start with overall churn rate
df_temp_churn = df['churn'].value_counts(normalize = True).reset_index()
df_temp_churn
```

```
[4]: churn  proportion
0      0      0.902848
1      1      0.097152
```

[5]: *#before going further we will see the churn percentages*

```
plt.bar('Churn', df_temp_churn.iloc[0,1], label = 'Retention')
plt.bar('Churn', df_temp_churn.iloc[1,1], bottom = df_temp_churn.iloc[0,1],
      ↪,label = 'Churn')
plt.ylabel('Percentage')
plt.title('Churn (Stacked Bar)')
plt.text('Churn', df_temp_churn.iloc[0,1]/2, f"{df_temp_churn.iloc[0,1]:.2%}",
      ↪ha='center', va='center', color='black', fontsize=12)
plt.text('Churn', df_temp_churn.iloc[0,1] + df_temp_churn.iloc[1,1]/2,
      ↪f"{df_temp_churn.iloc[1,1]:.2%}", ha='center', va='center', color='black',
      ↪fontsize=12)
plt.legend()
plt.show()
```




Overall, we have 9.72% churn rate which is more than what client expected. Our goal of this analysis is to study the dataset, observe the relationship between Churn and predictors, and make a predictive model for the client.

1.3.1 Date Analysis

```
[6]: #Analyzing dates data
df['tenure'] = (df['date_end'] - df['date_activ']).dt.days #creating tenure_
↳parameter
```

```
[7]: df['tenure'].hist()
plt.show()
```



```
[8]: def plot_distribution(dataframe, column, ax, bins_=50):
    """
    Plot variable distribution in a stacked histogram of churned or retained_
    ↳company
    """
    # Create a temporal dataframe with the data to be plot
    temp = pd.DataFrame({"Retention": dataframe[dataframe["churn"]==0][column],
                        "Churn":dataframe[dataframe["churn"]==1][column]})
    # Plot the histogram
    temp[["Retention","Churn"]].plot(kind='hist', bins=bins_, ax=ax,
    ↳stacked=True)
    # X-axis label
    ax.set_xlabel(column)
```


```
# Change the x-axis to plain style  
ax.ticklabel_format(style='plain', axis='x')
```

```
[9]: #Columns deemed not necessary for the task  
df.drop(columns = ['date_modif_prod', 'date_renewal'], inplace = True)
```

```
[10]: #creating ending month column to see if some months have more churn than others  
df['end_month'] = df['date_end'].dt.month  
df.drop(columns = ['date_activ', 'date_end'], inplace = True)
```

```
[11]: print(df.groupby('end_month').sum()['churn'])  
sns.lineplot(data = df.groupby('end_month').sum()['churn'].reset_index(), x =  
            ↪ 'end_month', y = 'churn')  
plt.ylim(bottom = 0)  
plt.show()
```

```
end_month  
1      112  
2      149  
3      124  
4       94  
5      108  
6      143  
7       89  
8      111  
9       98  
10     139  
11     143  
12     109  
Name: churn, dtype: int64
```



We are able to see three modes or peaks for churns happening across the years, February, June, and November. This signifies that while churn happens across the year it peaks at the start of the year, in the middle, and then in the end of the year with little variability across them.

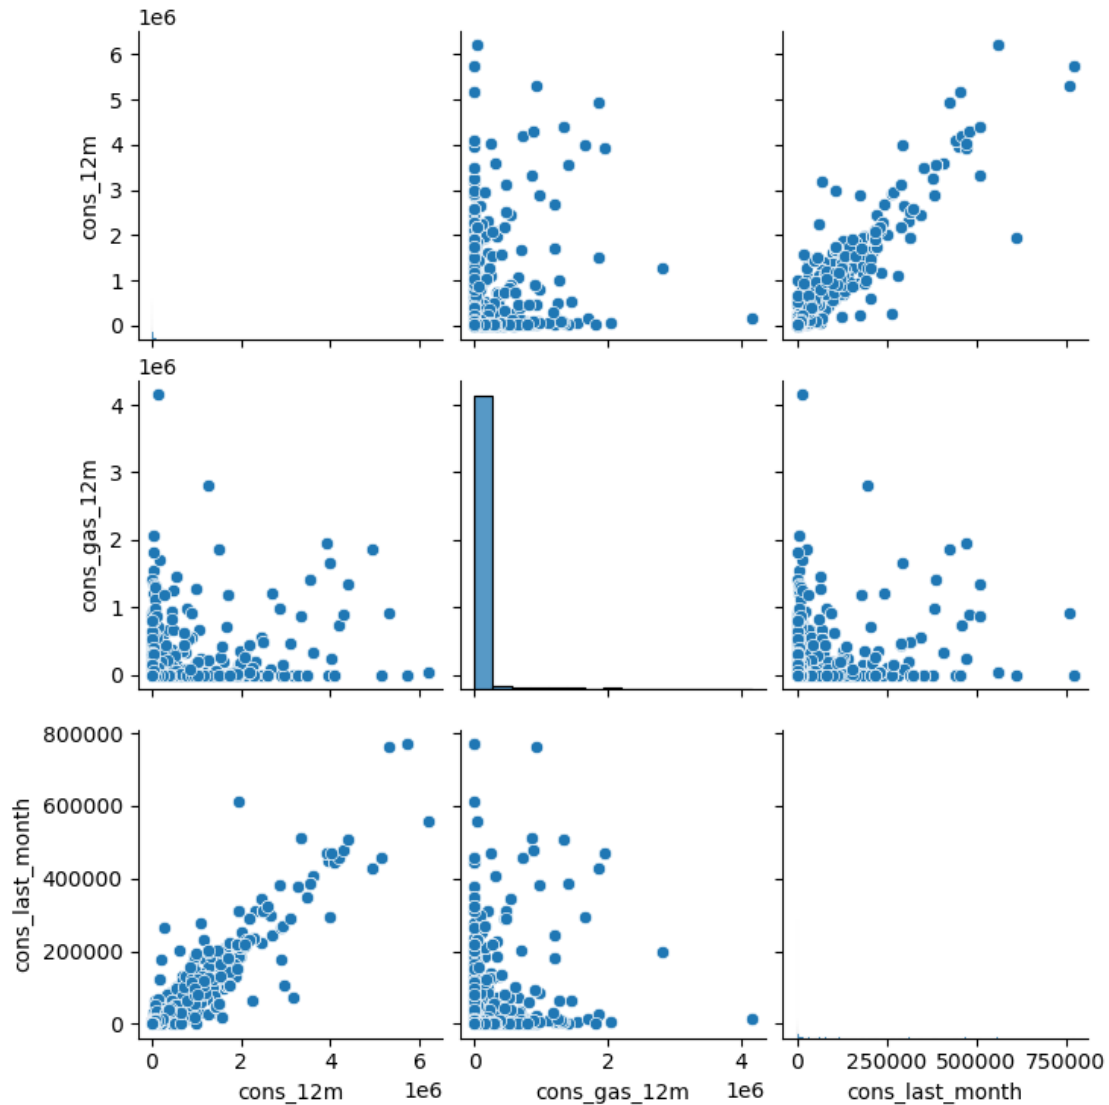
1.3.2 Consumption Variables

```
[12]: df.filter(regex = '^cons_').corr()
```

```
[12]:
```

	cons_12m	cons_gas_12m	cons_last_month
cons_12m	1.000000	0.488474	0.968212
cons_gas_12m	0.488474	1.000000	0.507007
cons_last_month	0.968212	0.507007	1.000000

```
[13]: sns.pairplot(df.filter(regex = '^cons_'))
plt.show()
```




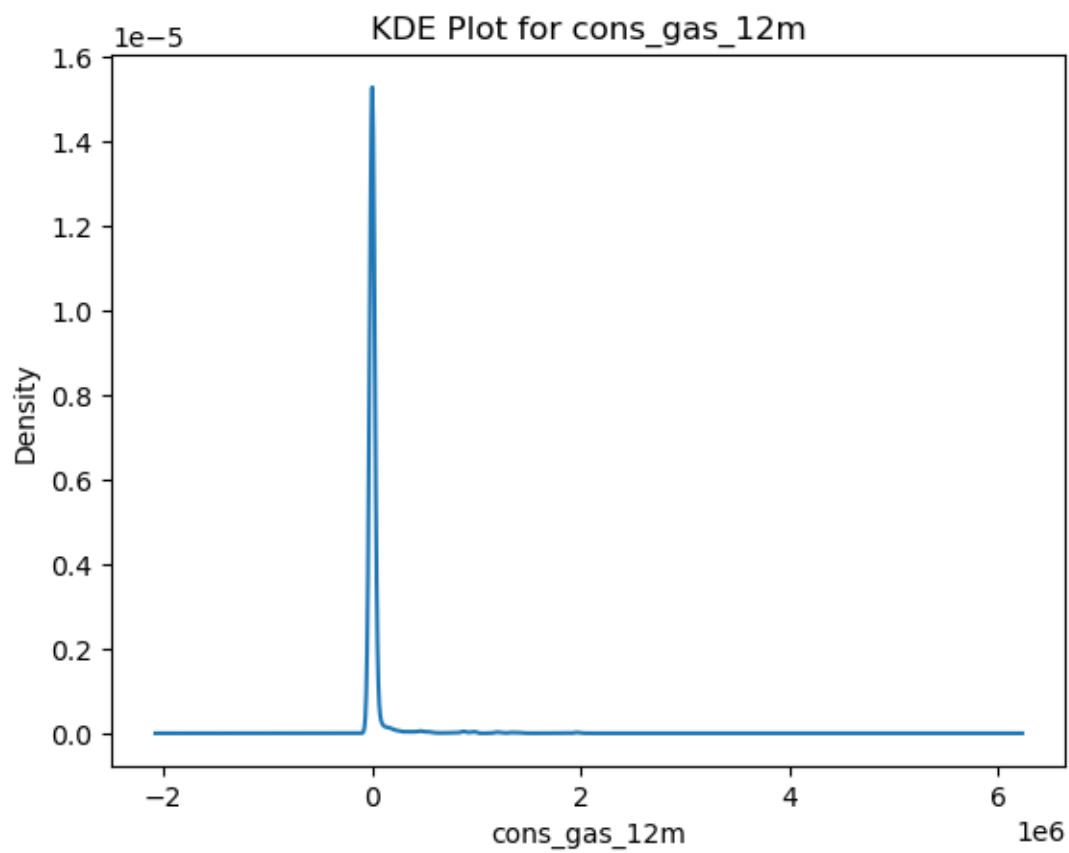
```
[14]: df.filter(regex = '^cons_').corr()
```

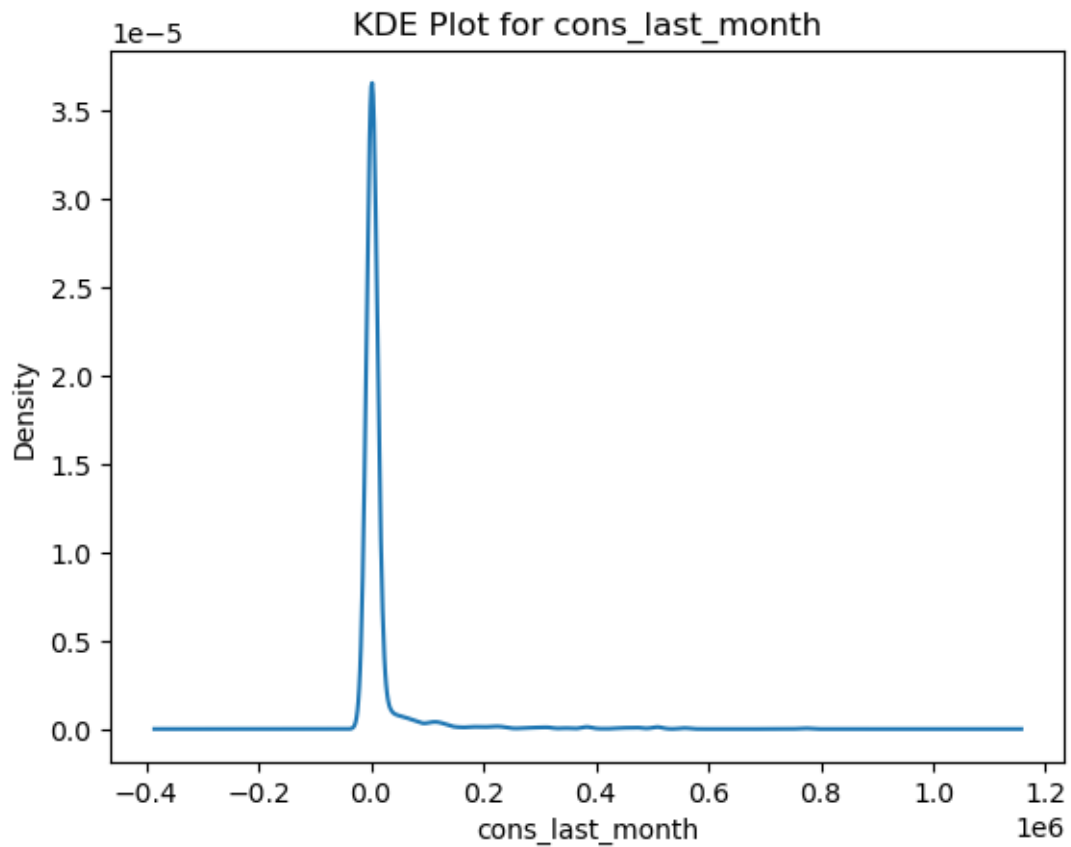
```
[14]:
```

	cons_12m	cons_gas_12m	cons_last_month
cons_12m	1.000000	0.488474	0.968212
cons_gas_12m	0.488474	1.000000	0.507007
cons_last_month	0.968212	0.507007	1.000000

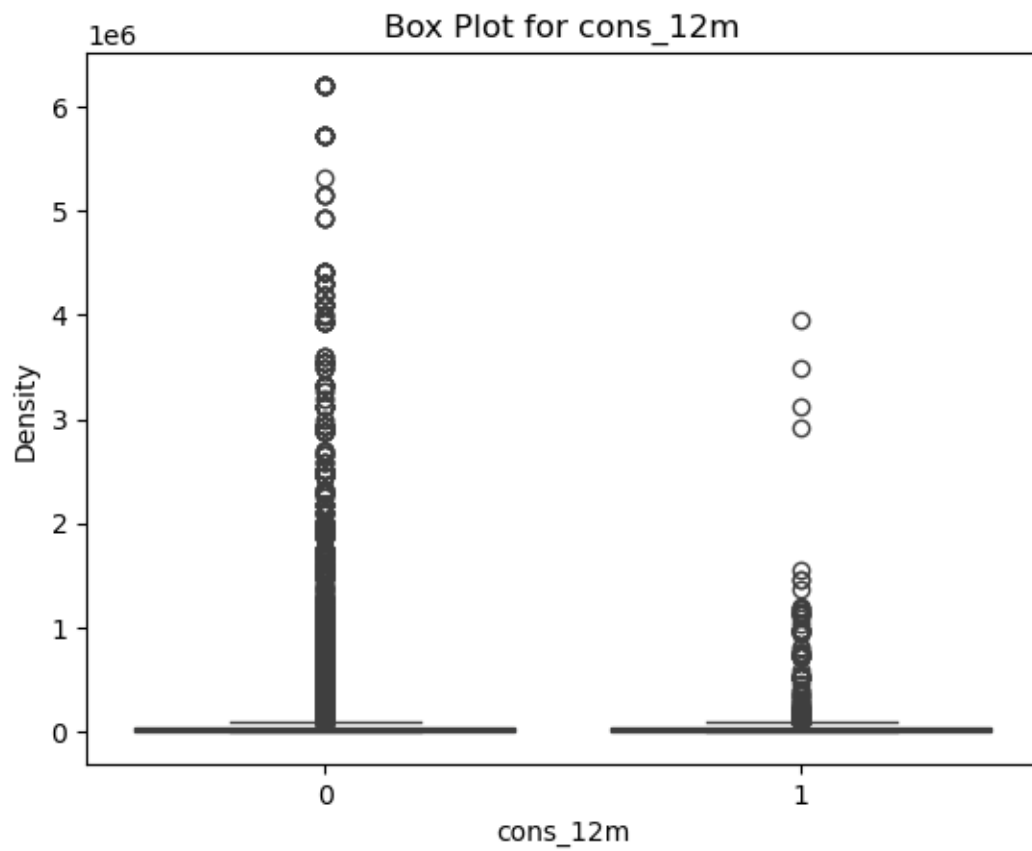
```
[15]: for col in df.filter(regex = '^cons_'):
df[col].plot.kde()
plt.title(f'KDE Plot for {col}')
plt.xlabel(col)
plt.ylabel('Density')
plt.show()
```





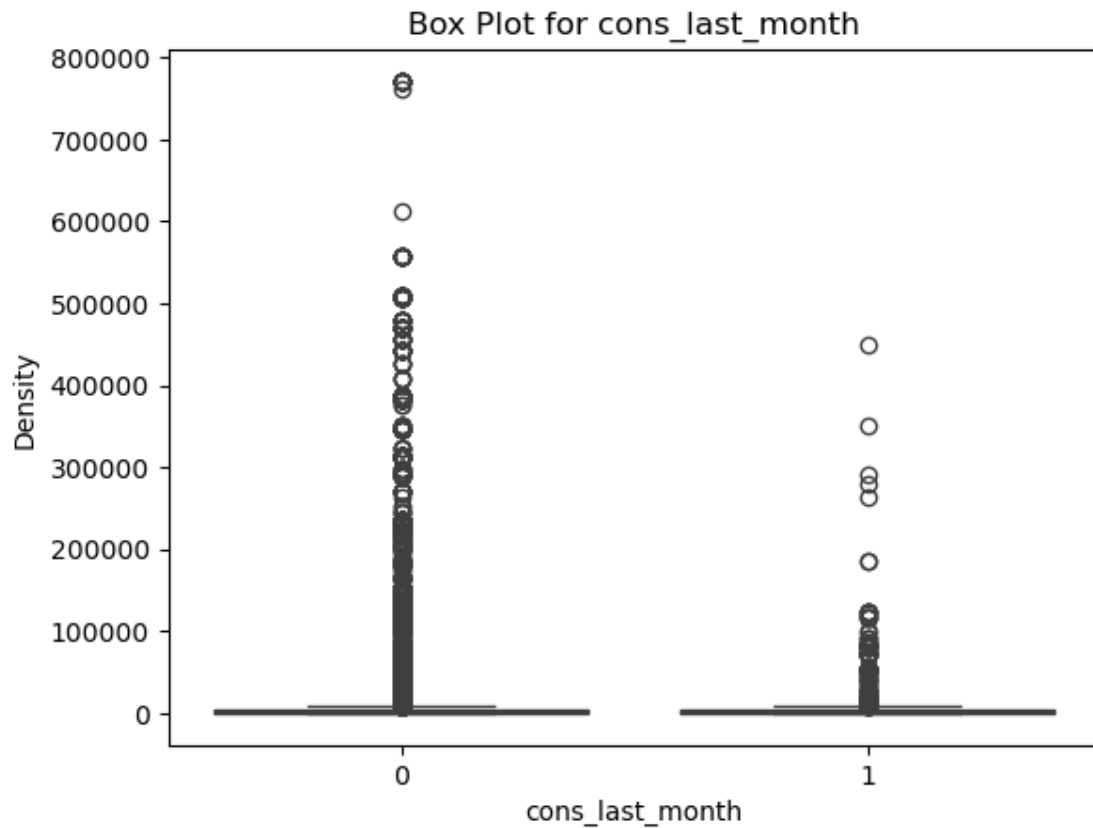




```
[16]: for col in df.filter(regex = '^cons_'):  
      plt.figure()  
      sns.boxplot(x = df['churn'], y = df[col])  
      plt.title(f'Box Plot for {col}')  
      plt.xlabel(col)  
      plt.ylabel('Density')  
      plt.show()
```











```
[17]: fig, axs = plt.subplots(nrows=4, figsize=(18, 25))

plot_distribution(df, 'cons_12m', axs[0])
plot_distribution(df[df['has_gas'] == 't'], 'cons_gas_12m', axs[1])
plot_distribution(df, 'cons_last_month', axs[2])
plot_distribution(df, 'imp_cons', axs[3])
plt.show()
```



```
[18]: df.filter(regex = '^cons_').describe()
```

```
[18]:
```

	cons_12m	cons_gas_12m	cons_last_month
count	1.460600e+04	1.460600e+04	14606.000000
mean	1.592203e+05	2.809238e+04	16090.269752
std	5.734653e+05	1.629731e+05	64364.196422
min	0.000000e+00	0.000000e+00	0.000000
25%	5.674750e+03	0.000000e+00	0.000000
50%	1.411550e+04	0.000000e+00	792.500000
75%	4.076375e+04	0.000000e+00	3383.000000
max	6.207104e+06	4.154590e+06	771203.000000

Consumption Variables are right skewed with most churns happening with customers with low consumption overall. Nevertheless, variables are highly skewed and there are a lot of outliers, and on top of that we are indeed in need for normalization for this dataset as range vary a lot across different categories.

1.3.3 Forecast

```
[19]: df.filter(regex = 'forecast_').describe()
```

```
[19]:
```

	forecast_cons_12m	forecast_cons_year	forecast_discount_energy \
count	14606.000000	14606.000000	14606.000000
mean	1868.614880	1399.762906	0.966726
std	2387.571531	3247.786255	5.108289
min	0.000000	0.000000	0.000000
25%	494.995000	0.000000	0.000000
50%	1112.875000	314.000000	0.000000
75%	2401.790000	1745.750000	0.000000
max	82902.830000	175375.000000	30.000000

	forecast_meter_rent_12m	forecast_price_energy_off_peak \
count	14606.000000	14606.000000
mean	63.086871	0.137283
std	66.165783	0.024623
min	0.000000	0.000000
25%	16.180000	0.116340
50%	18.795000	0.143166
75%	131.030000	0.146348
max	599.310000	0.273963

	forecast_price_energy_peak	forecast_price_pow_off_peak
count	14606.000000	14606.000000
mean	0.050491	43.130056
std	0.049037	4.485988
min	0.000000	0.000000
25%	0.000000	40.606701
50%	0.084138	44.311378
75%	0.098837	44.311378

max

0.195975

59.266378

```
[20]: df.filter(regex = 'forecast_').corr()
```

```
[20]:
```

	forecast_cons_12m	forecast_cons_year \
forecast_cons_12m	1.000000	0.647727
forecast_cons_year	0.647727	1.000000
forecast_discount_energy	0.058435	0.008518
forecast_meter_rent_12m	0.305627	0.276009
forecast_price_energy_off_peak	-0.135646	-0.158012
forecast_price_energy_peak	0.254056	0.251005
forecast_price_pow_off_peak	-0.018477	-0.044190

	forecast_discount_energy \
forecast_cons_12m	0.058435
forecast_cons_year	0.008518
forecast_discount_energy	1.000000
forecast_meter_rent_12m	-0.008388
forecast_price_energy_off_peak	0.353735
forecast_price_energy_peak	0.059318
forecast_price_pow_off_peak	0.048024

	forecast_meter_rent_12m \
forecast_cons_12m	0.305627
forecast_cons_year	0.276009
forecast_discount_energy	-0.008388
forecast_meter_rent_12m	1.000000
forecast_price_energy_off_peak	-0.579353
forecast_price_energy_peak	0.706376
forecast_price_pow_off_peak	-0.203089

	forecast_price_energy_off_peak \
forecast_cons_12m	-0.135646
forecast_cons_year	-0.158012
forecast_discount_energy	0.353735
forecast_meter_rent_12m	-0.579353
forecast_price_energy_off_peak	1.000000
forecast_price_energy_peak	-0.330138
forecast_price_pow_off_peak	0.630377

	forecast_price_energy_peak \
forecast_cons_12m	0.254056
forecast_cons_year	0.251005
forecast_discount_energy	0.059318
forecast_meter_rent_12m	0.706376
forecast_price_energy_off_peak	-0.330138
forecast_price_energy_peak	1.000000

```

forecast_price_pow_off_peak          -0.242017

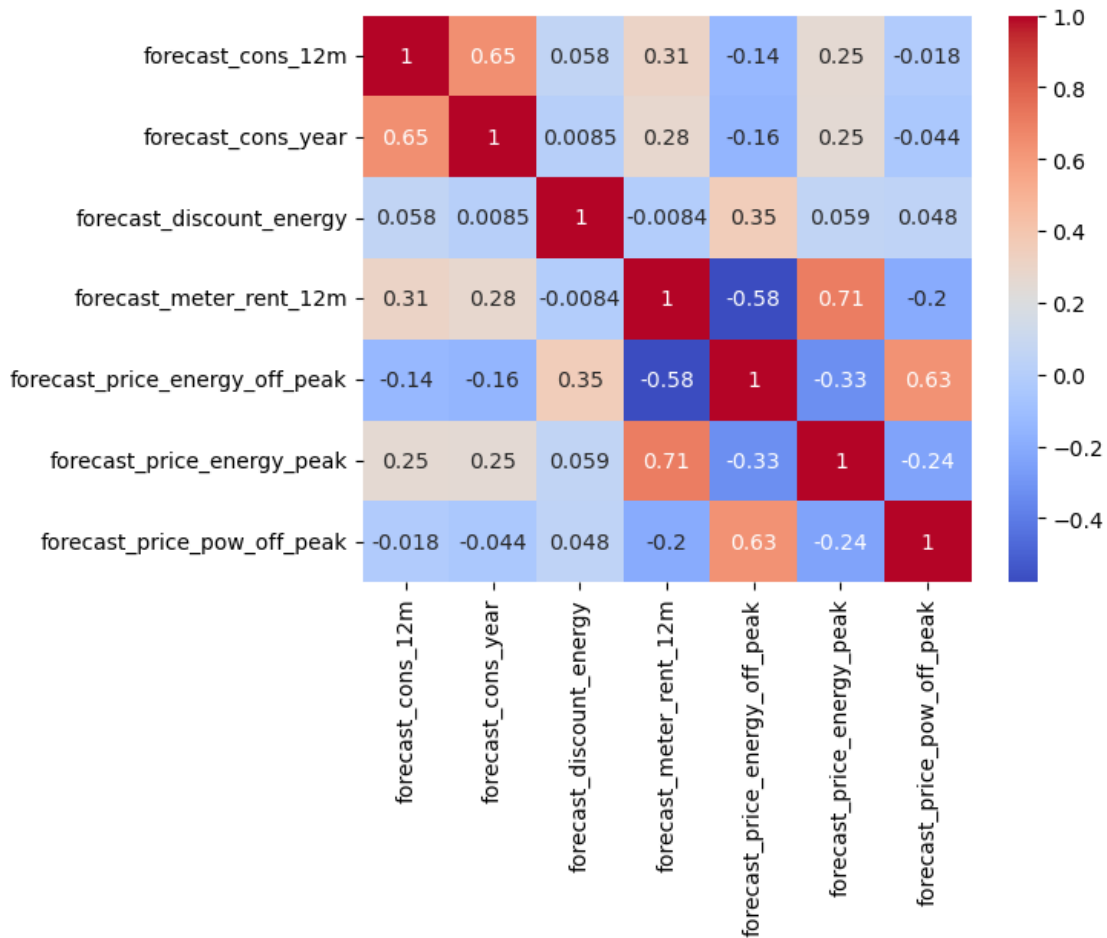
                                forecast_price_pow_off_peak
forecast_cons_12m                    -0.018477
forecast_cons_year                   -0.044190
forecast_discount_energy              0.048024
forecast_meter_rent_12m              -0.203089
forecast_price_energy_off_peak        0.630377
forecast_price_energy_peak            -0.242017
forecast_price_pow_off_peak           1.000000

```

```

[21]: sns.heatmap(df.filter(regex = 'forecast_').corr(), cmap = 'coolwarm', annot = True)
plt.show()

```

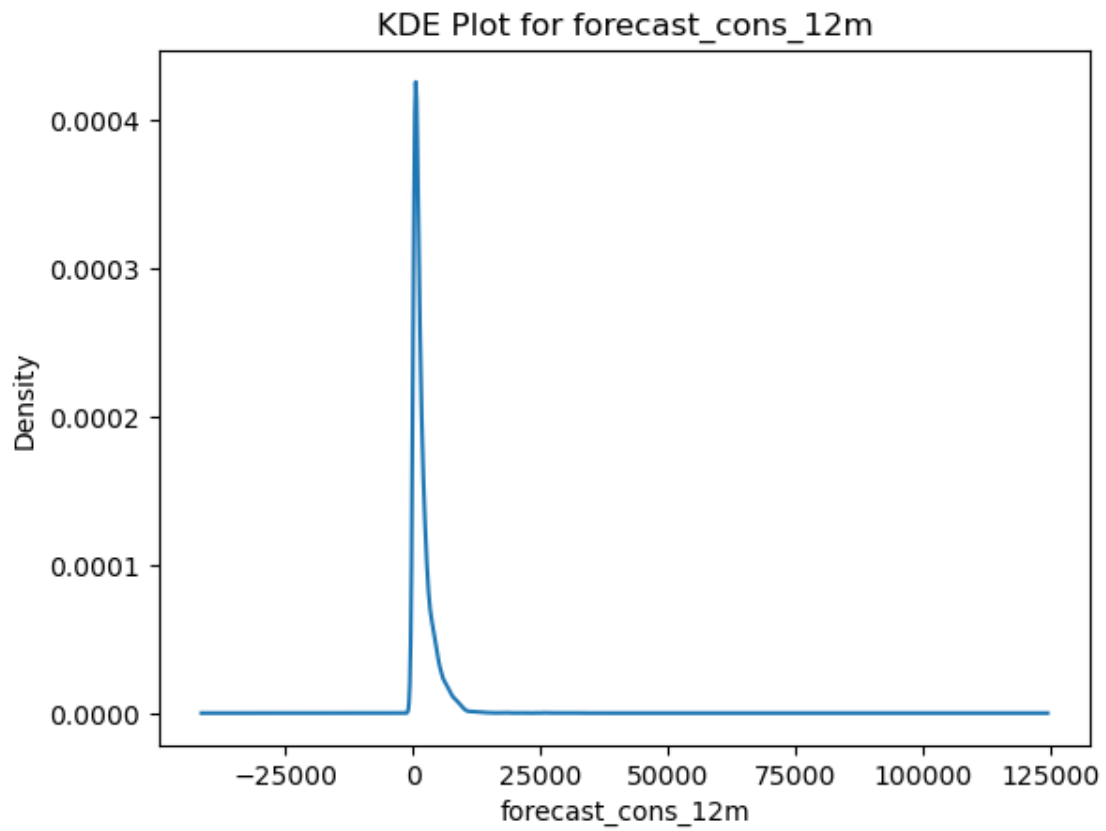


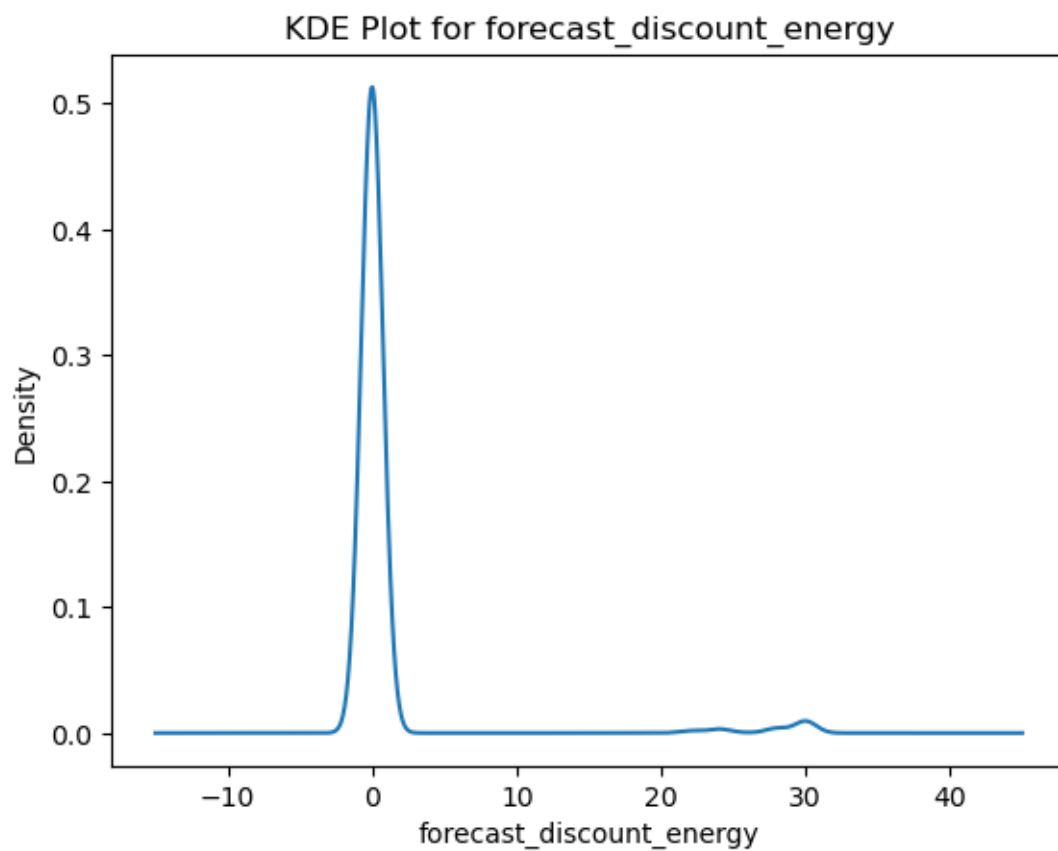
```

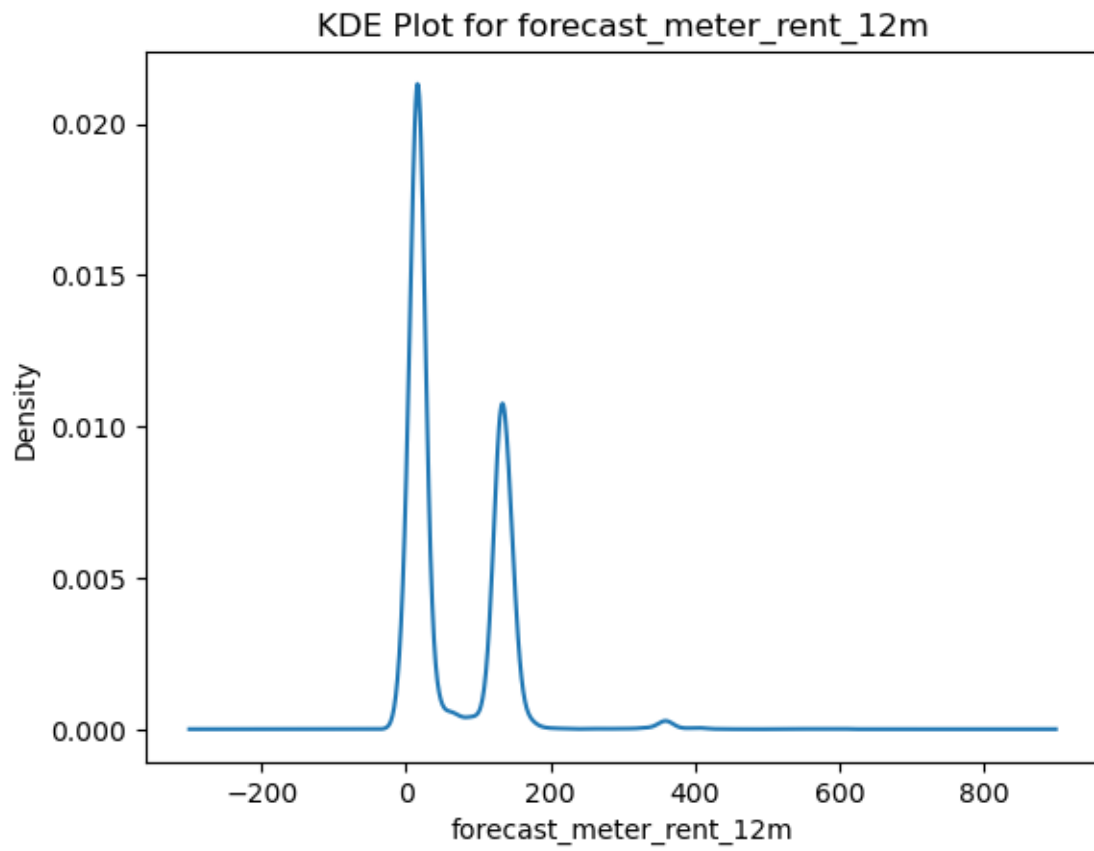
[22]: # Removing forecasted yearly consumption value as we will use 12 month forecast
df.drop(columns = 'forecast_cons_year', inplace = True)

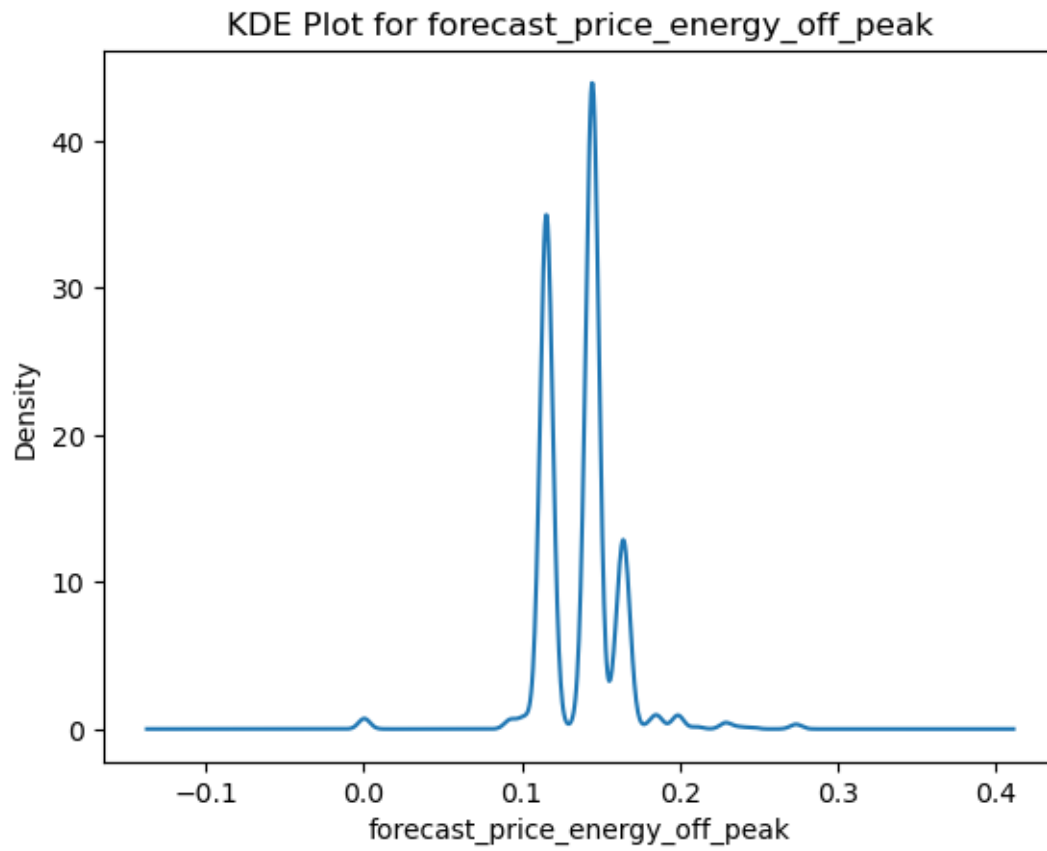
```

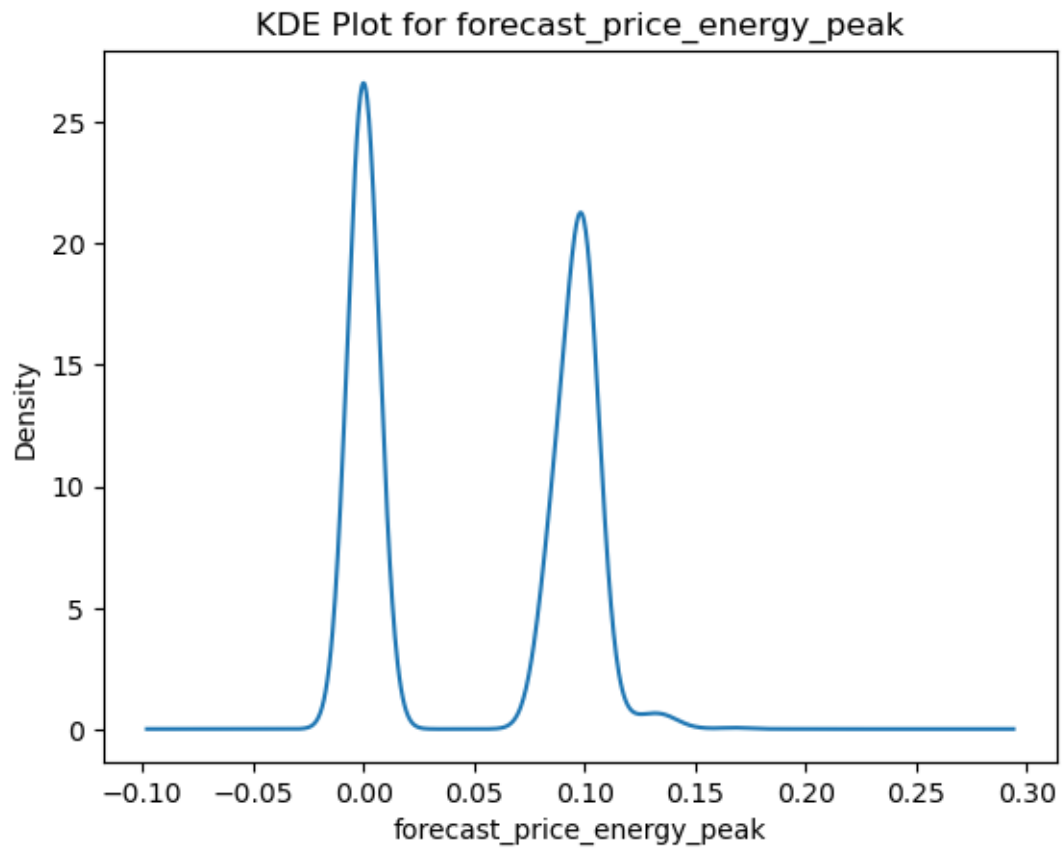
```
[23]: for col in df.filter(regex = '^forecast_'):  
      df[col].plot.kde()  
      plt.title(f'KDE Plot for {col}')  
      plt.xlabel(col)  
      plt.ylabel('Density')  
      plt.show()
```

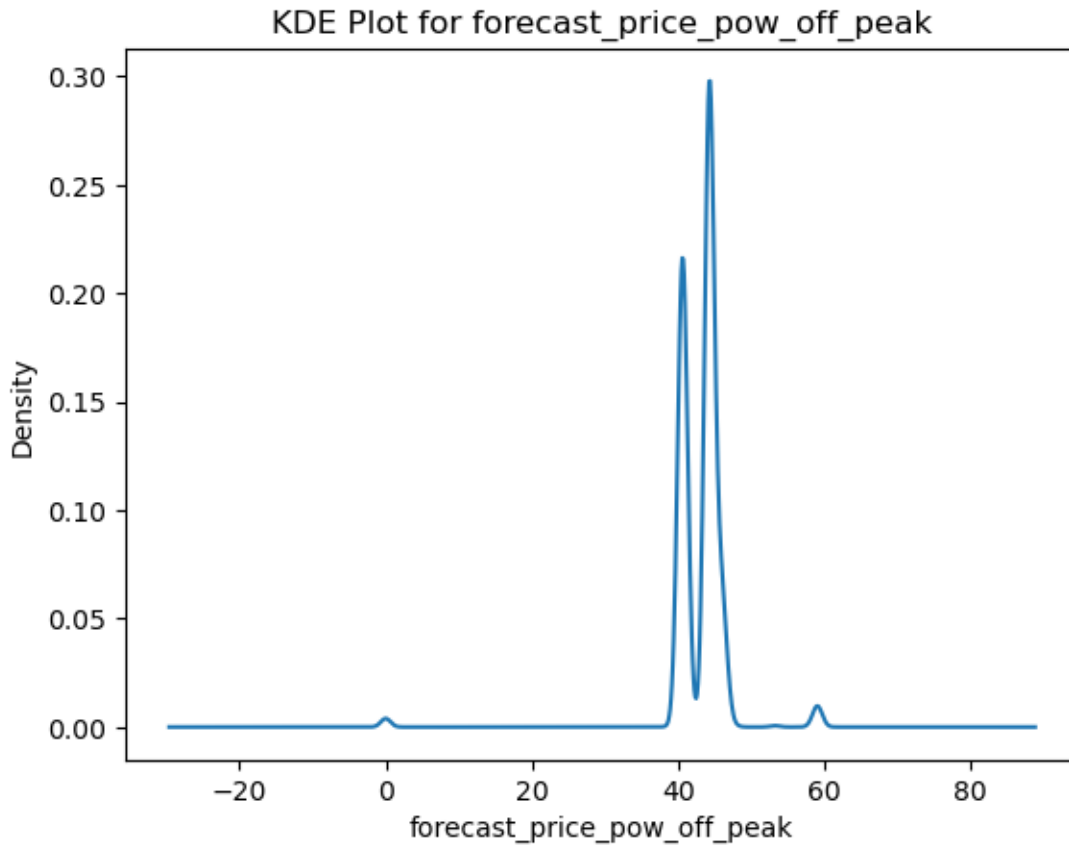








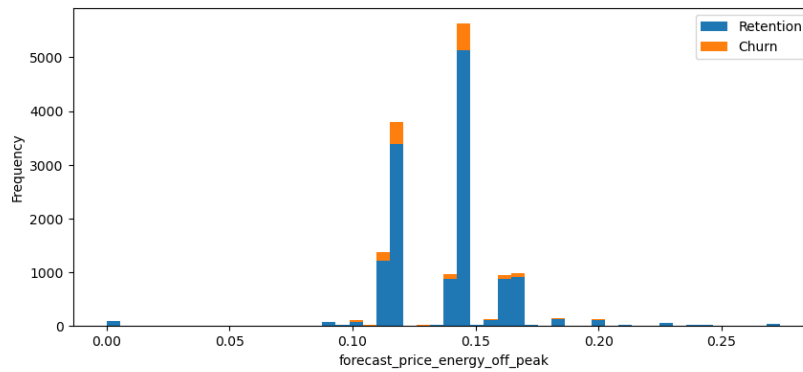
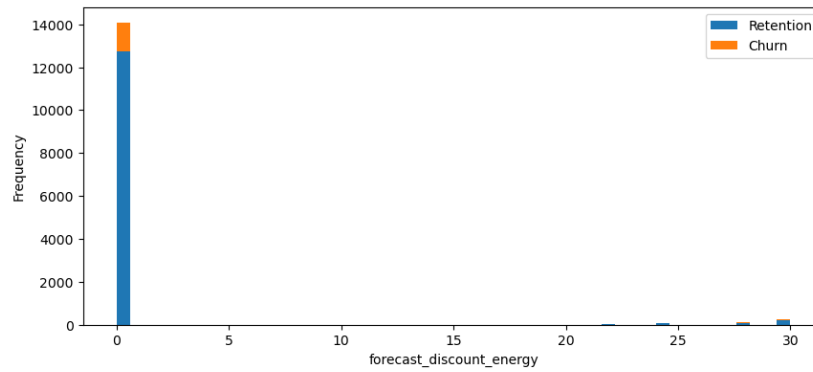
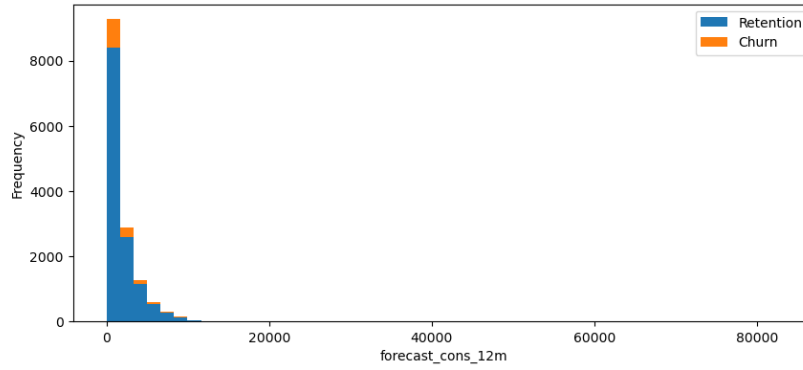






```
[24]: fig, axs = plt.subplots(nrows=5, figsize=(10, 25))

plot_distribution(df, 'forecast_cons_12m', axs[0])
plot_distribution(df, 'forecast_discount_energy', axs[1])
plot_distribution(df, 'forecast_price_energy_off_peak', axs[2])
plot_distribution(df, 'forecast_price_energy_peak', axs[3])
plot_distribution(df, 'forecast_meter_rent_12m', axs[4])
plt.show()
```

```
[25]: df.filter(regex = '^forecast_').describe()
```

```
[25]:
```

	forecast_cons_12m	forecast_discount_energy	forecast_meter_rent_12m	\
count	14606.000000	14606.000000	14606.000000	
mean	1868.614880	0.966726	63.086871	
std	2387.571531	5.108289	66.165783	
min	0.000000	0.000000	0.000000	
25%	494.995000	0.000000	16.180000	
50%	1112.875000	0.000000	18.795000	
75%	2401.790000	0.000000	131.030000	
max	82902.830000	30.000000	599.310000	

	forecast_price_energy_off_peak	forecast_price_energy_peak	\
count	14606.000000	14606.000000	
mean	0.137283	0.050491	
std	0.024623	0.049037	
min	0.000000	0.000000	
25%	0.116340	0.000000	
50%	0.143166	0.084138	
75%	0.146348	0.098837	
max	0.273963	0.195975	

	forecast_price_pow_off_peak
count	14606.000000
mean	43.130056
std	4.485988
min	0.000000
25%	40.606701
50%	44.311378
75%	44.311378
max	59.266378

Just like consumption data we have highly skewed variables for forecasting, which was expected. Our concern is why energy price for peak period is less than off peak periods. It should have been reversed, need confirmation if data was ported properly.

1.3.4 Others

Before analyzing other variables further it should be noted that variables covering variance of price over the period has been provided by BCGX without relevant mention in the datacard. We will cover them in brief besides other variables.

```
[ ]:
```

```
[26]: df['has_gas'] = df['has_gas'].replace(['t','f'], [1,0])
```

```
/tmp/ipykernel_153168/2844197130.py:1: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df['has_gas'] = df['has_gas'].replace(['t','f'], [1,0])
```

```
[27]: #we have around 18% of observations with gas connection
print(df.loc[df['has_gas'] == 1,:].shape[0]/df.shape[0])
```

```
0.18150075311515815
```

```
[28]: df.loc[df['has_gas'] == 1,:]['churn'].value_counts(normalize = True)
```

```
[28]: churn
0    0.918144
1    0.081856
Name: proportion, dtype: float64
```

```
[29]: df['churn'].value_counts(normalize = True)
```


```
[29]: churn
0    0.902848
1    0.097152
Name: proportion, dtype: float64
```


Churn rate for is quite lower for customers with gas connection.

```
[30]: #As expected gross and net margin are highly correlated. We will remove net_
margin for power subscription
print(df.filter(regex = 'margin').corr())
df.drop(columns = 'margin_net_pow_ele', inplace = True)
```

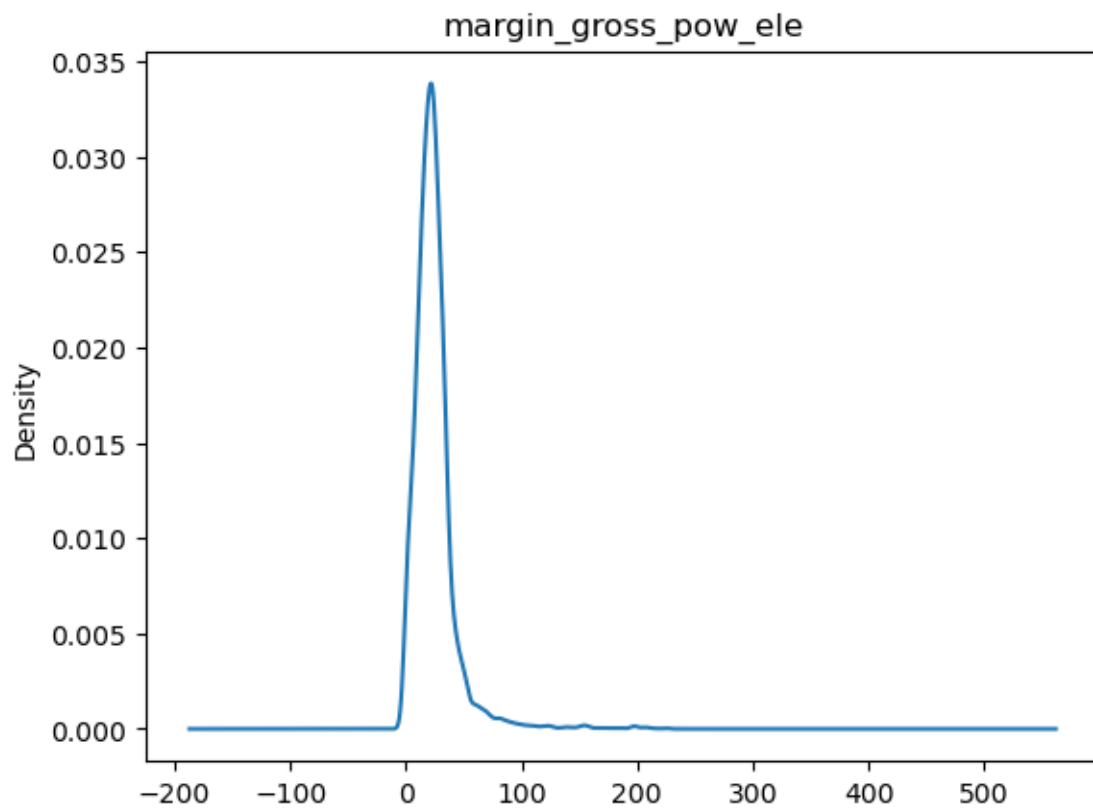
	margin_gross_pow_ele	margin_net_pow_ele	net_margin
margin_gross_pow_ele	1.000000	0.999914	0.031814
margin_net_pow_ele	0.999914	1.000000	0.031639
net_margin	0.031814	0.031639	1.000000

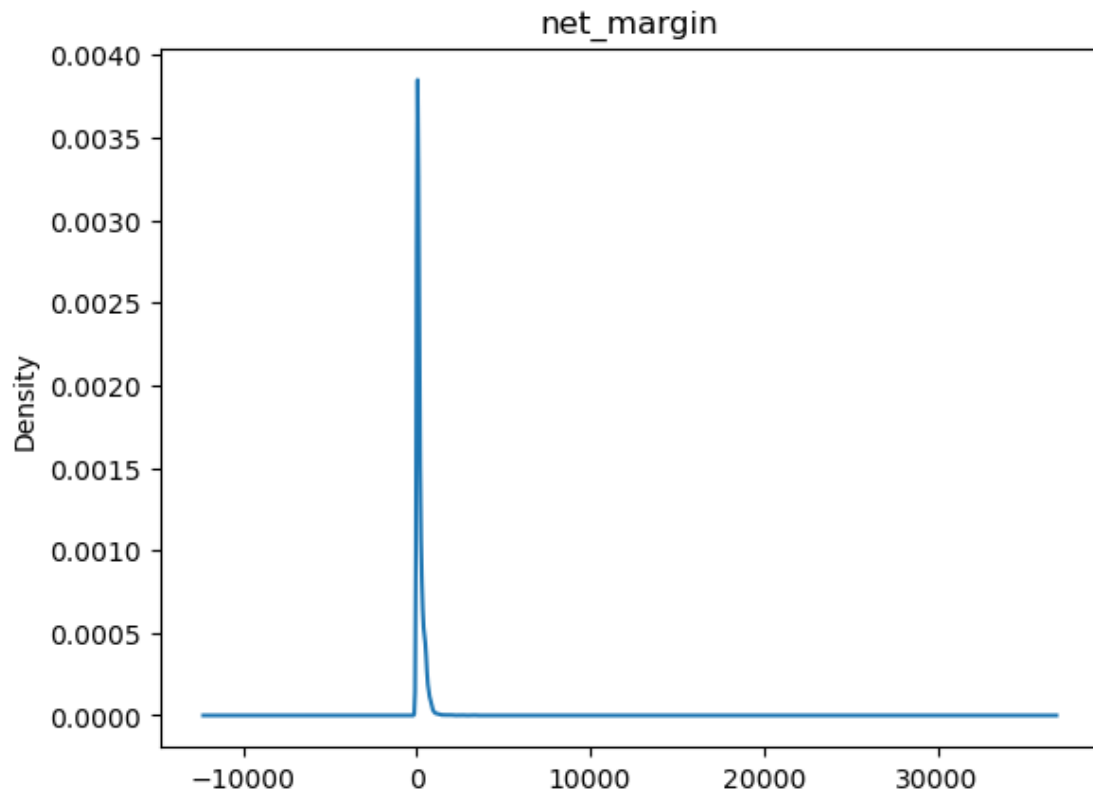
```
[31]: for cols in df.filter(regex = 'margin'):
sns.boxplot(x = df[cols])
plt.show()
```





```
[32]: for cols in df.filter(regex = 'margin'):  
      df[cols].plot.kde()  
      plt.title(f'{cols}')  
      plt.show()
```






```
[33]: df.filter(regex = 'margin').describe()
```

```
[33]:
```

	margin_gross_pow_ele	net_margin
count	14606.000000	14606.000000
mean	24.565121	189.264522
std	20.231172	311.798130
min	0.000000	0.000000
25%	14.280000	50.712500
50%	21.640000	112.530000
75%	29.880000	243.097500
max	374.640000	24570.650000

```
[34]: fig, axs = plt.subplots(nrows=2, figsize=(30, 25))

plot_distribution(df, 'margin_gross_pow_ele', axs[0])
plot_distribution(df, 'net_margin', axs[1])
plt.show()
```



Logically, net margin should be lower than gross margin but here net margin is higher. It maybe due to the fact that we have gross margin only for power subscription.

```
[35]: df.groupby('origin_up').sum()['churn'].sort_values(ascending = False)/df.
      ↪loc[df['churn'] == 1].shape[0]
```

```
[35]: origin_up
lxidpiddsbxsbosboudacockeimpuepw    0.629316
ldkssxwpmemidmecebumciepifcamkci    0.186047
kamkkxfxxuwbdslkwifmmcsiusuosws     0.181818
MISSING                               0.002819
ewxeelcelemmiwuafmddpobolfuxioce     0.000000
usapbepcfoloekilkwslwaxobdp          0.000000
Name: churn, dtype: float64
```

```
[36]: df.loc[df['churn'] == 0].groupby('origin_up').count()['id'].
      ↪sort_values(ascending = False)/df.loc[df['churn'] == 0].shape[0]
```



```
[36]: origin_up
      lxiidpiddsbxsbsoboudacockeimpuepw    0.470463
      kamkxxfxuwbdsldkwifmmsiusiusws     0.306059
      ldkssxwpmemidmecebumciepifcamkci    0.218700
      MISSING                             0.004550
      usapbepcfoloekilkwsdiboslwxobdp     0.000152
      ewxeelcelemmiwuafmddpobolfulxioce   0.000076
      Name: id, dtype: float64
```

This shows us that proportion of customers churning across the origin program varies. If it does not vary then proportion of churn vs not churn should have been same for the groups. But we can clearly observe that lxi Electricity Campaign has higher proportion of churn than other groups relatively.

```
[37]: df.groupby('channel_sales').sum()['churn'].sort_values(ascending = False)/df.
      ↪loc[df['churn'] == 1].shape[0]
```

```
[37]: channel_sales
      foosdfpfkusacimwkcsosbicdxkicaau    0.577872
      MISSING                             0.199436
      usilxuppasemublllopkaafesmlibmsdf   0.097252
      lmkebamcaaclubfxadlmueccxoimlema    0.072586
      ewpakwlliwisiwduibdlfmalxowmwpci    0.052854
      epumfxlbckeskwexbiuasklxalciuu      0.000000
      fixdbufsefwooaasfcxdxadsiekocaea    0.000000
      sddiedcslfslkckwlfkdpoeailfpeds     0.000000
      Name: churn, dtype: float64
```

```
[38]: df.loc[df['churn'] == 0].groupby('channel_sales').count()['id'].
      ↪sort_values(ascending = False)/df.loc[df['churn'] == 0].shape[0]
```

```
[38]: channel_sales
      foosdfpfkusacimwkcsosbicdxkicaau    0.449989
      MISSING                             0.261015
      lmkebamcaaclubfxadlmueccxoimlema    0.131948
      usilxuppasemublllopkaafesmlibmsdf   0.093805
      ewpakwlliwisiwduibdlfmalxowmwpci    0.062031
      sddiedcslfslkckwlfkdpoeailfpeds     0.000834
      epumfxlbckeskwexbiuasklxalciuu      0.000227
      fixdbufsefwooaasfcxdxadsiekocaea    0.000152
      Name: id, dtype: float64
```

```
[39]: print(df.groupby('channel_sales')['churn'].value_counts(normalize = True))
```

channel_sales	churn	
MISSING	0	0.924027
	1	0.075973
epumfxlbckeskwexbiuasklxalciuu	0	1.000000

ewpakwlliwisiwduibdlfmalxowmwpci	0	0.916013
	1	0.083987
fixdbufsefwooaasfcxdxadsiekoceaa	0	1.000000
foosdfpfkusacimwkcsosbicdxkicaua	0	0.878590
	1	0.121410
lmkebamcaaclubfxadlmueccxoimlema	0	0.944113
	1	0.055887
sddiedcslfslkckwlfdpoeailfpeds	0	1.000000
usilxuppasemublllopkaafesmlibmsdf	0	0.899636
	1	0.100364

Name: proportion, dtype: float64

```
[40]: #Analyzing the breakup more in depth
print(df.groupby('origin_up')['churn'].value_counts(normalize = True))
print(df.groupby(['channel_sales', 'origin_up']).agg({'churn': 'mean'}))
```

origin_up	churn
MISSING	0 0.937500
	1 0.062500
ewxeelcelemmiwuafmddpobolfuxioce	0 1.000000
kamkkxfxxuwbdslkwifmmcsiusiusws	0 0.939916
	1 0.060084
ldkssxwpmemidmecebumciepifcamkci	0 0.916137
	1 0.083863
lxidpiddsbxsbosboudacockeimpuepw	0 0.874172
	1 0.125828
usapbecpfoloekilkwsdiboslwaxobdp	0 1.000000

Name: proportion, dtype: float64


channel_sales	origin_up	churn
MISSING	MISSING	0.000000
	ewxeelcelemmiwuafmddpobolfuxioce	0.000000
	kamkkxfxxuwbdslkwifmmcsiusiusws	0.065330
	ldkssxwpmemidmecebumciepifcamkci	0.085601
	lxidpiddsbxsbosboudacockeimpuepw	0.080793
	usapbecpfoloekilkwsdiboslwaxobdp	0.000000
epumfxlbckeskwekxbiuasklxlalciiu	ldkssxwpmemidmecebumciepifcamkci	0.000000
	lxidpiddsbxsbosboudacockeimpuepw	0.000000
ewpakwlliwisiwduibdlfmalxowmwpci	MISSING	0.000000
	kamkkxfxxuwbdslkwifmmcsiusiusws	0.063452
	ldkssxwpmemidmecebumciepifcamkci	0.121107
	lxidpiddsbxsbosboudacockeimpuepw	0.073529
	usapbecpfoloekilkwsdiboslwaxobdp	0.000000
fixdbufsefwooaasfcxdxadsiekoceaa	kamkkxfxxuwbdslkwifmmcsiusiusws	0.000000
	ldkssxwpmemidmecebumciepifcamkci	0.000000
foosdfpfkusacimwkcsosbicdxkicaua	MISSING	0.093750
	kamkkxfxxuwbdslkwifmmcsiusiusws	0.067033
	ldkssxwpmemidmecebumciepifcamkci	0.084877

	lxidpiddsbxsbosboudacockeimpuepw	0.135747
lmkebamcaaclubfxadlmueccxoimlema	MISSING	0.142857
	kamkkxfixxuwbdslkwifmmcsiusiosws	0.036638
	ldkssxwpmemidmecebumciepifcamkci	0.054965
	lxidpiddsbxsbosboudacockeimpuepw	0.107558
sddiedcsflslkckwlfkdpoeetailfpeds	kamkkxfixxuwbdslkwifmmcsiusiosws	0.000000
	ldkssxwpmemidmecebumciepifcamkci	0.000000
	lxidpiddsbxsbosboudacockeimpuepw	0.000000
usilxuppasemublllopkaafesmlibmsdf	MISSING	0.000000
	kamkkxfixxuwbdslkwifmmcsiusiosws	0.071984
	ldkssxwpmemidmecebumciepifcamkci	0.108527
	lxidpiddsbxsbosboudacockeimpuepw	0.119835


For the sales channels and campaigns, proportion of customers is similar to the churn customers proportion but for some groups churn rate is quite higher. For electricity campaigns we can observe that lxidpi has 12.6% churn rate which is quite higher than 9.72% churn rate generally. And for sales channel its 12.12% for foosdf channel. Overall, some sales channel and campaigns witnessed more than average customer churn rate.

And combining both of them in subgroups we can see that there are many subgroups with more than 12% churn rate with highest being at 14.3% churn rate which is 47% higher than average churn rate. All this is a metter of concerns and imply that churn rate varies a lot across different sales channel and promotion campaigns.

```
[41]: sns.barplot(data = df, x = 'channel_sales', y = 'pow_max')
plt.xticks(rotation = 80)
plt.title('Subscribed Power across different Sales Channels')
plt.show()
```



```
[42]: sns.barplot(data = df, x = 'channel_sales', y = 'margin_gross_pow_ele')
plt.xticks(rotation = 80)
plt.title('Subscribed Power across different Sales Channels')
plt.show()
```



```
[43]: #Analyzing antiquity of the clients
df['num_years_antig'].describe()
```

```
[43]: count    14606.000000
      mean       4.997809
      std       1.611749
```

```

min          1.000000
25%          4.000000
50%          5.000000
75%          6.000000
max          13.000000
Name: num_years_antig, dtype: float64

```

```
[44]: df.groupby('channel_sales').count()['num_years_antig']
```

```

[44]: channel_sales
MISSING          3725
epumfxlbckeskwekxbiuasklxalciuu          3
ewpakwlliwisiwduibdlfmalxowmwpci        893
fixdbufsefwooaasfcxdxadsiekocaea          2
foosdfpfkusacimwkcsosbicdxkicaua       6754
lmkebamcaaclubfxadlmueccxoimlema       1843
sddiedcslfslkckwlfdpoeaailfpeds         11
usilxuppasemublllopkaafesmlibmsdf      1375
Name: num_years_antig, dtype: int64

```

```

[45]: temp_df = pd.merge(df.groupby('channel_sales').sum()['churn'].reset_index(),df.
    ↪groupby('channel_sales')[['num_years_antig']].mean().reset_index(), on =_
    ↪'channel_sales', how = 'inner')

temp_df.columns = ['channel_sales', 'churn', 'mean_years_antig']
temp_df

```

```

[45]:
      channel_sales  churn  mean_years_antig
0      MISSING      283         6.492617
1  epumfxlbckeskwekxbiuasklxalciuu          0         3.000000
2  ewpakwlliwisiwduibdlfmalxowmwpci         75         4.272116
3  fixdbufsefwooaasfcxdxadsiekocaea          0         3.000000
4  foosdfpfkusacimwkcsosbicdxkicaua        820         4.414125
5  lmkebamcaaclubfxadlmueccxoimlema        103         5.196419
6  sddiedcslfslkckwlfdpoeaailfpeds          0         3.000000
7  usilxuppasemublllopkaafesmlibmsdf        138         4.043636

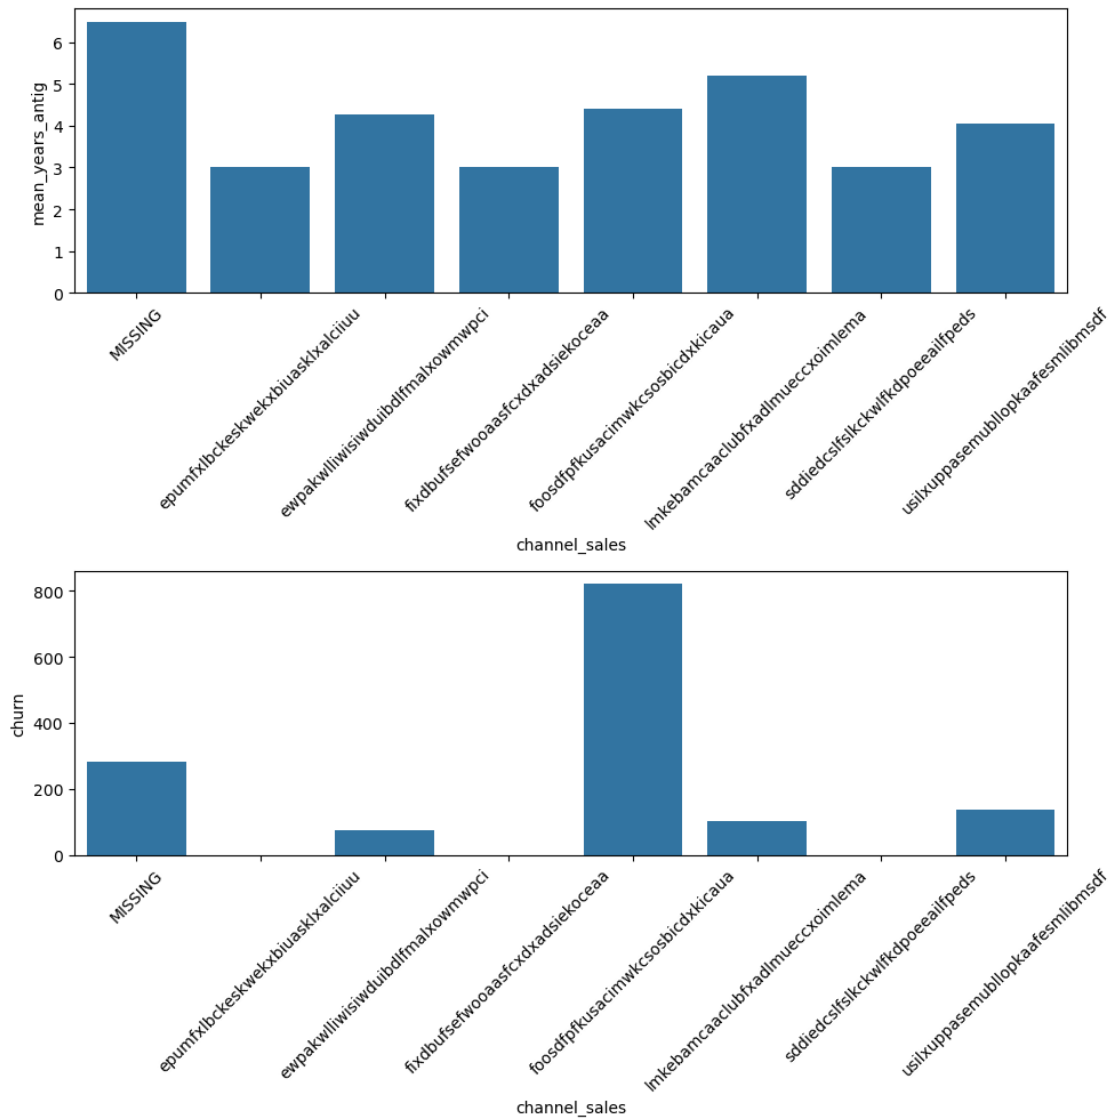
```

```


[46]: fig, ax = plt.subplots(nrows = 2, figsize = (10,10))

sns.barplot(temp_df, x = 'channel_sales', y = 'mean_years_antig', ax = ax[0])
sns.barplot(temp_df, x = 'channel_sales', y = 'churn', ax = ax[1])
for ax in ax:
    ax.tick_params(axis='x', labelrotation=45)
plt.tight_layout()
plt.show()

```



```
[47]: sns.boxplot(df, x = 'churn', y = 'num_years_antig')
plt.show()
```



Overall, customers churned has mean of atinquity is around 4 years. While distribution of churn and retention for customers is quite similar when we look across different sales channel it presents a different story.

We can observe churns happening in sales channels having antiquity more than the mean of 4 years. While the churn rate appears normal for all groups, ‘foosdfpfkusacimwkcsoibcdxkicaua’ sales group has a higher amount of churns than others. Yes, it has more customers too but even then it should have been 500-600 churns instead of 820 churns. This implies that this sales group witnessed higher than normal churn rate for PowerCo.

```
[48]: df.filter(regex = 'var_').corr()
```

```
[48]:
```

	var_year_price_off_peak_var \
var_year_price_off_peak_var	1.000000
var_year_price_peak_var	0.268477
var_year_price_mid_peak_var	0.245311
var_year_price_off_peak_fix	0.515049
var_year_price_peak_fix	0.227269
var_year_price_mid_peak_fix	0.229893
var_year_price_off_peak	0.515064
var_year_price_peak	0.227272
var_year_price_mid_peak	0.229894

var_6m_price_off_peak_var	0.619343
var_6m_price_peak_var	0.265968
var_6m_price_mid_peak_var	0.269555
var_6m_price_off_peak_fix	0.441878
var_6m_price_peak_fix	0.264435
var_6m_price_mid_peak_fix	0.264312
var_6m_price_off_peak	0.441880
var_6m_price_peak	0.264437
var_6m_price_mid_peak	0.264313

	var_year_price_peak_var \
var_year_price_off_peak_var	0.268477
var_year_price_peak_var	1.000000
var_year_price_mid_peak_var	0.515645
var_year_price_off_peak_fix	0.298254
var_year_price_peak_fix	0.451123
var_year_price_mid_peak_fix	0.475760
var_year_price_off_peak	0.298256
var_year_price_peak	0.451135
var_year_price_mid_peak	0.475761
var_6m_price_off_peak_var	0.295638
var_6m_price_peak_var	0.647083
var_6m_price_mid_peak_var	0.312539
var_6m_price_off_peak_fix	0.278371
var_6m_price_peak_fix	0.279135
var_6m_price_mid_peak_fix	0.283905
var_6m_price_off_peak	0.278372
var_6m_price_peak	0.279144
var_6m_price_mid_peak	0.283906

	var_year_price_mid_peak_var \
var_year_price_off_peak_var	0.245311
var_year_price_peak_var	0.515645
var_year_price_mid_peak_var	1.000000
var_year_price_off_peak_fix	0.296689
var_year_price_peak_fix	0.948578
var_year_price_mid_peak_fix	0.961481
var_year_price_off_peak	0.296691
var_year_price_peak	0.948579
var_year_price_mid_peak	0.961483
var_6m_price_off_peak_var	0.301787
var_6m_price_peak_var	0.292145
var_6m_price_mid_peak_var	0.666421
var_6m_price_off_peak_fix	0.276626
var_6m_price_peak_fix	0.637463
var_6m_price_mid_peak_fix	0.638773
var_6m_price_off_peak	0.276627

var_6m_price_peak	0.637464
var_6m_price_mid_peak	0.638774

	var_year_price_off_peak_fix \
var_year_price_off_peak_var	0.515049
var_year_price_peak_var	0.298254
var_year_price_mid_peak_var	0.296689
var_year_price_off_peak_fix	1.000000
var_year_price_peak_fix	0.340722
var_year_price_mid_peak_fix	0.310471
var_year_price_off_peak	1.000000
var_year_price_peak	0.340724
var_year_price_mid_peak	0.310471
var_6m_price_off_peak_var	0.755875
var_6m_price_peak_var	0.425467
var_6m_price_mid_peak_var	0.420256
var_6m_price_off_peak_fix	0.919000
var_6m_price_peak_fix	0.460190
var_6m_price_mid_peak_fix	0.426243
var_6m_price_off_peak	0.919000
var_6m_price_peak	0.460194
var_6m_price_mid_peak	0.426243

	var_year_price_peak_fix \
var_year_price_off_peak_var	0.227269
var_year_price_peak_var	0.451123
var_year_price_mid_peak_var	0.948578
var_year_price_off_peak_fix	0.340722
var_year_price_peak_fix	1.000000
var_year_price_mid_peak_fix	0.954353
var_year_price_off_peak	0.340723
var_year_price_peak	1.000000
var_year_price_mid_peak	0.954354
var_6m_price_off_peak_var	0.309477
var_6m_price_peak_var	0.273434
var_6m_price_mid_peak_var	0.649140
var_6m_price_off_peak_fix	0.310610
var_6m_price_peak_fix	0.662291
var_6m_price_mid_peak_fix	0.637355
var_6m_price_off_peak	0.310611
var_6m_price_peak	0.662291
var_6m_price_mid_peak	0.637355

	var_year_price_mid_peak_fix \
var_year_price_off_peak_var	0.229893
var_year_price_peak_var	0.475760
var_year_price_mid_peak_var	0.961481

var_year_price_off_peak_fix	0.310471
var_year_price_peak_fix	0.954353
var_year_price_mid_peak_fix	1.000000
var_year_price_off_peak	0.310472
var_year_price_peak	0.954354
var_year_price_mid_peak	1.000000
var_6m_price_off_peak_var	0.310865
var_6m_price_peak_var	0.283188
var_6m_price_mid_peak_var	0.664660
var_6m_price_off_peak_fix	0.290756
var_6m_price_peak_fix	0.650680
var_6m_price_mid_peak_fix	0.670265
var_6m_price_off_peak	0.290757
var_6m_price_peak	0.650680
var_6m_price_mid_peak	0.670265

	var_year_price_off_peak	var_year_price_peak \
var_year_price_off_peak_var	0.515064	0.227272
var_year_price_peak_var	0.298256	0.451135
var_year_price_mid_peak_var	0.296691	0.948579
var_year_price_off_peak_fix	1.000000	0.340724
var_year_price_peak_fix	0.340723	1.000000
var_year_price_mid_peak_fix	0.310472	0.954354
var_year_price_off_peak	1.000000	0.340725
var_year_price_peak	0.340725	1.000000
var_year_price_mid_peak	0.310472	0.954355
var_6m_price_off_peak_var	0.755879	0.309480
var_6m_price_peak_var	0.425468	0.273442
var_6m_price_mid_peak_var	0.420257	0.649141
var_6m_price_off_peak_fix	0.918999	0.310612
var_6m_price_peak_fix	0.460191	0.662291
var_6m_price_mid_peak_fix	0.426243	0.637355
var_6m_price_off_peak	0.918999	0.310613
var_6m_price_peak	0.460195	0.662291
var_6m_price_mid_peak	0.426244	0.637355

	var_year_price_mid_peak \
var_year_price_off_peak_var	0.229894
var_year_price_peak_var	0.475761
var_year_price_mid_peak_var	0.961483
var_year_price_off_peak_fix	0.310471
var_year_price_peak_fix	0.954354
var_year_price_mid_peak_fix	1.000000
var_year_price_off_peak	0.310472
var_year_price_peak	0.954355
var_year_price_mid_peak	1.000000
var_6m_price_off_peak_var	0.310865

var_6m_price_peak_var	0.283188
var_6m_price_mid_peak_var	0.664661
var_6m_price_off_peak_fix	0.290756
var_6m_price_peak_fix	0.650680
var_6m_price_mid_peak_fix	0.670265
var_6m_price_off_peak	0.290757
var_6m_price_peak	0.650681
var_6m_price_mid_peak	0.670265

	var_6m_price_off_peak_var	var_6m_price_peak_var \
var_year_price_off_peak_var	0.619343	0.265968
var_year_price_peak_var	0.295638	0.647083
var_year_price_mid_peak_var	0.301787	0.292145
var_year_price_off_peak_fix	0.755875	0.425467
var_year_price_peak_fix	0.309477	0.273434
var_year_price_mid_peak_fix	0.310865	0.283188
var_year_price_off_peak	0.755879	0.425468
var_year_price_peak	0.309480	0.273442
var_year_price_mid_peak	0.310865	0.283188
var_6m_price_off_peak_var	1.000000	0.468700
var_6m_price_peak_var	0.468700	1.000000
var_6m_price_mid_peak_var	0.475458	0.418850
var_6m_price_off_peak_fix	0.813497	0.459064
var_6m_price_peak_fix	0.480090	0.398631
var_6m_price_mid_peak_fix	0.473619	0.403081
var_6m_price_off_peak	0.813501	0.459065
var_6m_price_peak	0.480094	0.398645
var_6m_price_mid_peak	0.473620	0.403081

	var_6m_price_mid_peak_var \
var_year_price_off_peak_var	0.269555
var_year_price_peak_var	0.312539
var_year_price_mid_peak_var	0.666421
var_year_price_off_peak_fix	0.420256
var_year_price_peak_fix	0.649140
var_year_price_mid_peak_fix	0.664660
var_year_price_off_peak	0.420257
var_year_price_peak	0.649141
var_year_price_mid_peak	0.664661
var_6m_price_off_peak_var	0.475458
var_6m_price_peak_var	0.418850
var_6m_price_mid_peak_var	1.000000
var_6m_price_off_peak_fix	0.450765
var_6m_price_peak_fix	0.970093
var_6m_price_mid_peak_fix	0.976743
var_6m_price_off_peak	0.450766
var_6m_price_peak	0.970093

var_6m_price_mid_peak	0.976744	
-----------------------	----------	--


	var_6m_price_off_peak_fix	var_6m_price_peak_fix \
var_year_price_off_peak_var	0.441878	0.264435
var_year_price_peak_var	0.278371	0.279135
var_year_price_mid_peak_var	0.276626	0.637463
var_year_price_off_peak_fix	0.919000	0.460190
var_year_price_peak_fix	0.310610	0.662291
var_year_price_mid_peak_fix	0.290756	0.650680
var_year_price_off_peak	0.918999	0.460191
var_year_price_peak	0.310612	0.662291
var_year_price_mid_peak	0.290756	0.650680
var_6m_price_off_peak_var	0.813497	0.480090
var_6m_price_peak_var	0.459064	0.398631
var_6m_price_mid_peak_var	0.450765	0.970093
var_6m_price_off_peak_fix	1.000000	0.495431
var_6m_price_peak_fix	0.495431	1.000000
var_6m_price_mid_peak_fix	0.460428	0.970213
var_6m_price_off_peak	1.000000	0.495432
var_6m_price_peak	0.495436	1.000000
var_6m_price_mid_peak	0.460428	0.970214

	var_6m_price_mid_peak_fix	var_6m_price_off_peak \
var_year_price_off_peak_var	0.264312	0.441880
var_year_price_peak_var	0.283905	0.278372
var_year_price_mid_peak_var	0.638773	0.276627
var_year_price_off_peak_fix	0.426243	0.919000
var_year_price_peak_fix	0.637355	0.310611
var_year_price_mid_peak_fix	0.670265	0.290757
var_year_price_off_peak	0.426243	0.918999
var_year_price_peak	0.637355	0.310613
var_year_price_mid_peak	0.670265	0.290757
var_6m_price_off_peak_var	0.473619	0.813501
var_6m_price_peak_var	0.403081	0.459065
var_6m_price_mid_peak_var	0.976743	0.450766
var_6m_price_off_peak_fix	0.460428	1.000000
var_6m_price_peak_fix	0.970213	0.495432
var_6m_price_mid_peak_fix	1.000000	0.460429
var_6m_price_off_peak	0.460429	1.000000
var_6m_price_peak	0.970214	0.495436
var_6m_price_mid_peak	1.000000	0.460429

	var_6m_price_peak	var_6m_price_mid_peak
var_year_price_off_peak_var	0.264437	0.264313
var_year_price_peak_var	0.279144	0.283906
var_year_price_mid_peak_var	0.637464	0.638774
var_year_price_off_peak_fix	0.460194	0.426243

var_year_price_peak_fix	0.662291	0.637355
var_year_price_mid_peak_fix	0.650680	0.670265
var_year_price_off_peak	0.460195	0.426244
var_year_price_peak	0.662291	0.637355
var_year_price_mid_peak	0.650681	0.670265
var_6m_price_off_peak_var	0.480094	0.473620
var_6m_price_peak_var	0.398645	0.403081
var_6m_price_mid_peak_var	0.970093	0.976744
var_6m_price_off_peak_fix	0.495436	0.460428
var_6m_price_peak_fix	1.000000	0.970214
var_6m_price_mid_peak_fix	0.970214	1.000000
var_6m_price_off_peak	0.495436	0.460429
var_6m_price_peak	1.000000	0.970214
var_6m_price_mid_peak	0.970214	1.000000

```
[49]: sns.heatmap(df.filter(regex = 'var').corr(), cmap = 'coolwarm')
plt.show()
```





These are the columns added by BCGX team, we can observe high correlations across many of the columns. While we can remove some of them, for prediction we will use decision trees and neural networks so this will not impact a lot.

```
[50]: df.groupby('churn').sum()[['imp_cons']]
```

```
[50]:      imp_cons
churn
0      2017138.69
1      214466.71
```

```
[51]: fig, axs = plt.subplots(nrows=2, figsize=(20, 20))

plot_distribution(df, 'imp_cons', axs[0])
plot_distribution(df, 'pow_max', axs[1])
plt.show()
```



```
[52]: df.groupby('origin_up').sum()['churn']
```

```
[52]: origin_up
MISSING                                4
ewxeelcelemmiwuafmddpobolfuxioce      0
kamkkxfxxuwbdslkwifmmcsiusiusws      258
ldkssxwpmemidmecebumciepifcamkci      264
lxidpiddsbxsbosboudacockeimpuepw      893
usapbecpfoloekilkwsdiboslwaxobdp       0
Name: churn, dtype: int64
```

1.3.5 Price Dataset EDA

```
[53]: #Loading price data
df_price = pd.read_csv('price_data.csv')
df_price['price_date'] = pd.to_datetime(df_price['price_date'])
df_price['price_month'] = df_price['price_date'].dt.month #Converting date_
↳column and extracting month
df_price.head()
```

```
[53]:
```

	id	price_date	price_off_peak_var	\
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	
3	038af19179925da21a25619c5a24b745	2015-04-01	0.149626	
4	038af19179925da21a25619c5a24b745	2015-05-01	0.149626	

	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	\
0	0.0	0.0	44.266931	0.0	
1	0.0	0.0	44.266931	0.0	
2	0.0	0.0	44.266931	0.0	
3	0.0	0.0	44.266931	0.0	
4	0.0	0.0	44.266931	0.0	

	price_mid_peak_fix	price_month
0	0.0	1
1	0.0	2
2	0.0	3
3	0.0	4
4	0.0	5

```
[54]: #extracting meaningful insights from date
```



```
df_price['is_weekend'] = df_price['price_date'].dt.dayofweek >= 5 #Weekend
    ↪price is generally more than weekdays
df_price['is_month_start'] = df_price['price_date'].dt.is_month_start #Month
    ↪start price
df_price['is_month_end'] = df_price['price_date'].dt.is_month_end #Month end
    ↪price
df_price['quarter'] = df_price['price_date'].dt.quarter #Quarterly price
```

```
[55]: df_price.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    193002 non-null object
1   price_date            193002 non-null datetime64[ns]
2   price_off_peak_var    193002 non-null float64
3   price_peak_var        193002 non-null float64
4   price_mid_peak_var    193002 non-null float64
5   price_off_peak_fix    193002 non-null float64
6   price_peak_fix        193002 non-null float64
7   price_mid_peak_fix    193002 non-null float64
8   price_month           193002 non-null int32
9   is_weekend            193002 non-null bool
10  is_month_start        193002 non-null bool
11  is_month_end          193002 non-null bool
12  quarter               193002 non-null int32
dtypes: bool(3), datetime64[ns](1), float64(6), int32(2), object(1)
memory usage: 13.8+ MB
```

```
[56]: #More feature engineering on price variables. We are covering price
    ↪differences, average prices, and weekend prices
```

```
df_price['var_peak_offpeak_spread'] = df_price['price_peak_var'] -
    ↪df_price['price_off_peak_var']
df_price['fix_peak_offpeak_spread'] = df_price['price_peak_fix'] -
    ↪df_price['price_off_peak_fix']
df_price['avg_var_price'] = df_price[['price_off_peak_var', 'price_peak_var',
    ↪'price_mid_peak_var']].mean(axis=1)
df_price['avg_fix_price'] = df_price[['price_off_peak_fix', 'price_peak_fix',
    ↪'price_mid_peak_fix']].mean(axis=1)
df_price['peak_var_x_is_weekend'] = df_price['price_peak_var'] *
    ↪df_price['is_weekend'].astype(int)
```

```
[57]: df_price.corr(numeric_only = True)
```

[57]:

	price_off_peak_var	price_peak_var \
price_off_peak_var	1.000000	-0.328580
price_peak_var	-0.328580	1.000000
price_mid_peak_var	-0.594872	0.821353
price_off_peak_fix	0.417097	-0.098627
price_peak_fix	-0.630018	0.796097
price_mid_peak_fix	-0.572229	0.807759
price_month	-0.083796	-0.012409
is_weekend	-0.016939	-0.003991
is_month_start	NaN	NaN
is_month_end	NaN	NaN
quarter	-0.081647	-0.011674
var_peak_offpeak_spread	-0.660111	0.926359
fix_peak_offpeak_spread	-0.742420	0.771820
avg_var_price	-0.174635	0.963302
avg_fix_price	-0.502655	0.780654
peak_var_x_is_weekend	-0.155150	0.427270

	price_mid_peak_var	price_off_peak_fix \
price_off_peak_var	-0.594872	0.417097
price_peak_var	0.821353	-0.098627
price_mid_peak_var	1.000000	-0.137848
price_off_peak_fix	-0.137848	1.000000
price_peak_fix	0.973960	-0.000428
price_mid_peak_fix	0.979717	-0.252661
price_month	0.003895	0.013119
is_weekend	0.001135	-0.001878
is_month_start	NaN	NaN
is_month_end	NaN	NaN
quarter	0.004327	0.012386
var_peak_offpeak_spread	0.890470	-0.244772
fix_peak_offpeak_spread	0.950930	-0.388586
avg_var_price	0.840718	0.006916
avg_fix_price	0.947287	0.168202
peak_var_x_is_weekend	0.355659	-0.044383

	price_peak_fix	price_mid_peak_fix	price_month \
price_off_peak_var	-0.630018	-0.572229	-0.083796
price_peak_var	0.796097	0.807759	-0.012409
price_mid_peak_var	0.973960	0.979717	0.003895
price_off_peak_fix	-0.000428	-0.252661	0.013119
price_peak_fix	1.000000	0.927308	-0.000862
price_mid_peak_fix	0.927308	1.000000	-0.001079
price_month	-0.000862	-0.001079	1.000000
is_weekend	0.001788	0.002100	-0.102431
is_month_start	NaN	NaN	NaN
is_month_end	NaN	NaN	NaN

quarter	-0.000344	-0.000579	0.971625
var_peak_offpeak_spread	0.884398	0.870628	0.023547
fix_peak_offpeak_spread	0.921579	0.952514	-0.005887
avg_var_price	0.799150	0.829294	-0.034696
avg_fix_price	0.980064	0.895341	0.002519
peak_var_x_is_weekend	0.345639	0.350978	-0.074579

	is_weekend	is_month_start	is_month_end	quarter \
price_off_peak_var	-0.016939	NaN	NaN	-0.081647
price_peak_var	-0.003991	NaN	NaN	-0.011674
price_mid_peak_var	0.001135	NaN	NaN	0.004327
price_off_peak_fix	-0.001878	NaN	NaN	0.012386
price_peak_fix	0.001788	NaN	NaN	-0.000344
price_mid_peak_fix	0.002100	NaN	NaN	-0.000579
price_month	-0.102431	NaN	NaN	0.971625
is_weekend	1.000000	NaN	NaN	-0.158085
is_month_start	NaN	NaN	NaN	NaN
is_month_end	NaN	NaN	NaN	NaN
quarter	-0.158085	NaN	NaN	1.000000
var_peak_offpeak_spread	0.003581	NaN	NaN	0.023275
fix_peak_offpeak_spread	0.002377	NaN	NaN	-0.005125
avg_var_price	-0.007841	NaN	NaN	-0.033266
avg_fix_price	0.001424	NaN	NaN	0.002840
peak_var_x_is_weekend	0.665309	NaN	NaN	-0.112290

	var_peak_offpeak_spread	fix_peak_offpeak_spread \
price_off_peak_var	-0.660111	-0.742420
price_peak_var	0.926359	0.771820
price_mid_peak_var	0.890470	0.950930
price_off_peak_fix	-0.244772	-0.388586
price_peak_fix	0.884398	0.921579
price_mid_peak_fix	0.870628	0.952514
price_month	0.023547	-0.005887
is_weekend	0.003581	0.002377
is_month_start	NaN	NaN
is_month_end	NaN	NaN
quarter	0.023275	-0.005125
var_peak_offpeak_spread	1.000000	0.909914
fix_peak_offpeak_spread	0.909914	1.000000
avg_var_price	0.835781	0.733662
avg_fix_price	0.821326	0.837749
peak_var_x_is_weekend	0.401691	0.335706

	avg_var_price	avg_fix_price	peak_var_x_is_weekend
price_off_peak_var	-0.174635	-0.502655	-0.155150
price_peak_var	0.963302	0.780654	0.427270
price_mid_peak_var	0.840718	0.947287	0.355659

price_off_peak_fix	0.006916	0.168202	-0.044383
price_peak_fix	0.799150	0.980064	0.345639
price_mid_peak_fix	0.829294	0.895341	0.350978
price_month	-0.034696	0.002519	-0.074579
is_weekend	-0.007841	0.001424	0.665309
is_month_start	NaN	NaN	NaN
is_month_end	NaN	NaN	NaN
quarter	-0.033266	0.002840	-0.112290
var_peak_offpeak_spread	0.835781	0.821326	0.401691
fix_peak_offpeak_spread	0.733662	0.837749	0.335706
avg_var_price	1.000000	0.818669	0.408921
avg_fix_price	0.818669	1.000000	0.338626
peak_var_x_is_weekend	0.408921	0.338626	1.000000

As expected correlation is high for various columns. While multicollinearity is a negative aspect for a dataset decision trees, especially boosting techniques can counter it alongside Neural Network approaches.

```
[58]: df_price['price_month']
```

```
[58]: 0      1
      1      2
      2      3
      3      4
      4      5
      ..
192997    8
192998    9
192999   10
193000   11
193001   12
Name: price_month, Length: 193002, dtype: int32
```

```
[59]: df_price[['price_peak_var', 'price_month']].groupby('price_month').
      ↪agg({'price_peak_var': 'mean'})
```

```
[59]:           price_peak_var
price_month
1           0.054950
2           0.055053
3           0.055118
4           0.056035
5           0.055359
6           0.055255
7           0.055369
8           0.053605
9           0.053532
```

10	0.053713
11	0.053620
12	0.053957

```
[60]: df_price[['price_peak_fix', 'price_month']].groupby('price_month').
      ↪agg({'price_peak_fix': 'mean'})
```

```
[60]:
```

	price_peak_fix
price_month	
1	10.692921
2	10.673719
3	10.644489
4	10.647277
5	10.602453
6	10.415769
7	10.642236
8	10.661678
9	10.602979
10	10.605431
11	10.641489
12	10.644109

Average peak price for Energy and Power is quite similar across the different months for an year. Overall prices are not varying much across the year.

```
[61]: #Checking null values
      print(df_price.isnull().sum())
```

id	0
price_date	0
price_off_peak_var	0
price_peak_var	0
price_mid_peak_var	0
price_off_peak_fix	0
price_peak_fix	0
price_mid_peak_fix	0
price_month	0
is_weekend	0
is_month_start	0
is_month_end	0
quarter	0
var_peak_offpeak_spread	0
fix_peak_offpeak_spread	0
avg_var_price	0
avg_fix_price	0
peak_var_x_is_weekend	0
dtype: int64	

```
[62]: print(df_price.describe())
```

	price_date	price_off_peak_var	price_peak_var \
count	193002	193002.000000	193002.000000
mean	2015-06-16 12:50:49.933161216	0.141027	0.054630
min	2015-01-01 00:00:00	0.000000	0.000000
25%	2015-04-01 00:00:00	0.125976	0.000000
50%	2015-07-01 00:00:00	0.146033	0.085483
75%	2015-10-01 00:00:00	0.151635	0.101673
max	2015-12-01 00:00:00	0.280700	0.229788
std	NaN	0.025032	0.049924

	price_mid_peak_var	price_off_peak_fix	price_peak_fix \
count	193002.000000	193002.000000	193002.000000
mean	0.030496	43.334477	10.622875
min	0.000000	0.000000	0.000000
25%	0.000000	40.728885	0.000000
50%	0.000000	44.266930	0.000000
75%	0.072558	44.444710	24.339581
max	0.114102	59.444710	36.490692
std	0.036298	5.410297	12.841895

	price_mid_peak_fix	price_month	quarter \
count	193002.000000	193002.000000	193002.000000
mean	6.409984	6.501161	2.500368
min	0.000000	1.000000	1.000000
25%	0.000000	4.000000	2.000000
50%	0.000000	7.000000	3.000000
75%	16.226389	10.000000	4.000000
max	17.458221	12.000000	4.000000
std	7.773592	3.451935	1.117995

	var_peak_offpeak_spread	fix_peak_offpeak_spread	avg_var_price \
count	193002.000000	193002.000000	193002.000000
mean	-0.086397	-32.711602	0.075384
min	-0.280700	-59.444710	0.000000
25%	-0.148477	-44.266931	0.049852
50%	-0.082931	-44.266930	0.084289
75%	-0.022136	-16.291555	0.098433
max	0.096908	0.000000	0.163193
std	0.062772	13.937180	0.024746

	avg_fix_price	peak_var_x_is_weekend
count	193002.000000	193002.000000
mean	20.122446	0.018117
min	0.000000	0.000000
25%	14.755644	0.000000

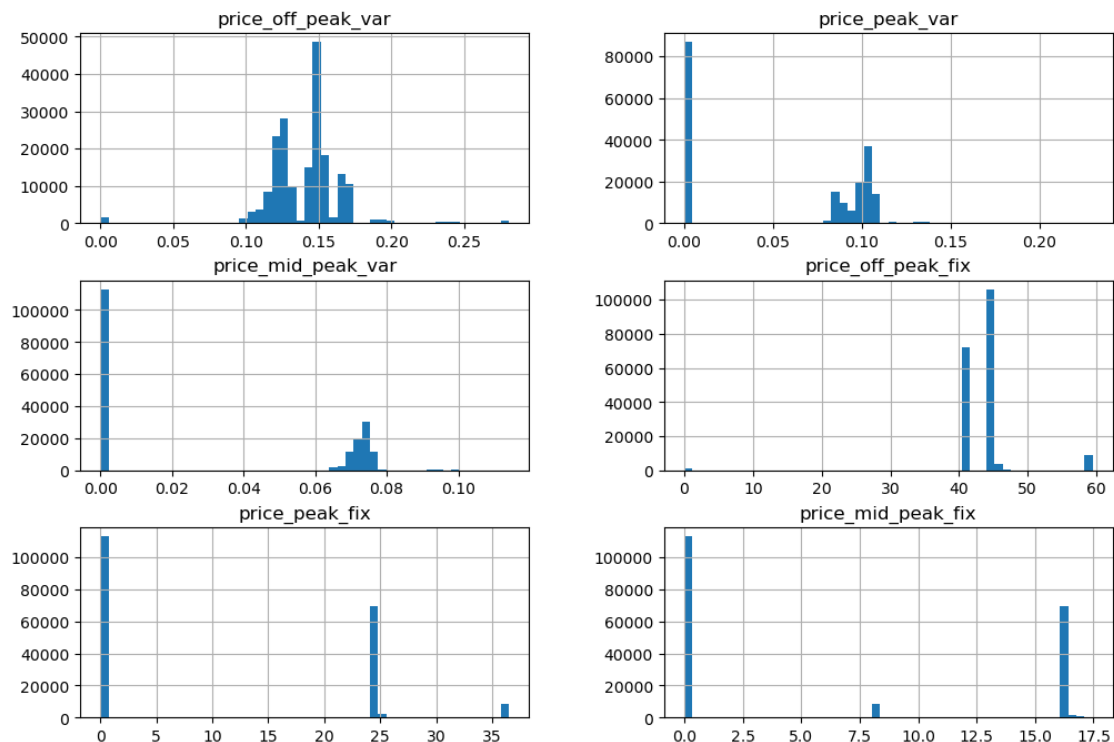
50%	14.814903	0.000000
75%	27.043980	0.000000
max	34.677296	0.229788
std	6.818637	0.038509

```
[63]: print(df_price[['is_weekend', 'is_month_start', 'is_month_end']].sum())
print(df_price['quarter'].value_counts())
```


```
is_weekend      64337
is_month_start  193002
is_month_end      0
dtype: int64
quarter
4      48266
3      48266
2      48244
1      48226
Name: count, dtype: int64
```

```
[64]: df_price[['price_off_peak_var', 'price_peak_var', 'price_mid_peak_var',
               'price_off_peak_fix', 'price_peak_fix', 'price_mid_peak_fix']].
       hist(bins=50, figsize=(12,8))
plt.suptitle('Distribution of Price Columns')
plt.show()
```

Distribution of Price Columns




```
[65]: df_price[['price_off_peak_var', 'price_peak_var', 'price_mid_peak_var']].
      plot(kind='box', figsize=(8,5))
plt.title('Boxplot of Variable Prices')
plt.show()
```



There are a lot of outliers in off peak vs other categories. And again we are seeing peak pricing less than off peak when in reality it happens in inverse, most probably BCGX team adjusted the variables for the sake of anonymity of their client. Regardless of name of variables we can use the data for our predictive tasks.

```
[66]: daily_avg = df_price.groupby('price_date')[['avg_var_price', 'avg_fix_price']].
      mean()
daily_avg.plot(figsize=(12,6), title='Average Energy(var) and Power(fix) Price_
      Over Time')
plt.ylabel('Price')
plt.show()
```





```
[67]: #Average price for energy and power does not change much across the year
daily_avg.head()
```


```
[67]:
```

price_date	avg_var_price	avg_fix_price
2015-01-01	0.075950	20.125141
2015-02-01	0.076063	20.120299
2015-03-01	0.076206	20.109548
2015-04-01	0.076595	20.119624
2015-05-01	0.076429	20.102882

```
[68]: df_price.groupby('price_month')['avg_var_price'].mean().plot(kind='bar',
    ↳title='Avg Energy Price by Month')
plt.ylabel('Avg Var Price')
plt.show()
```




```
[69]: df_price.groupby('price_month')['avg_fix_price'].mean().plot(kind='bar',  
    title='Avg Power Price by Month')  
plt.ylabel('Avg Fix Price')  
plt.show()
```




Average price for both energy and power connections are not varying much, client PowerCo has very stable pricing for the period.

```
[70]: df_price.groupby('is_weekend')['avg_var_price'].mean().plot(kind='bar',  
      title='Avg Energy Price: Weekend vs Weekday')  
plt.ylabel('Avg Var Price')  
plt.show()
```



```
[71]: df_price.groupby('is_weekend')['avg_fix_price'].mean().plot(kind='bar',  
      title='Avg Power Price: Weekend vs Weekday')  
plt.ylabel('Avg Fix Price')  
plt.show()
```



```
[72]: print(df_price.groupby('is_weekend')['avg_var_price'].mean())
```

```
is_weekend
False      0.075522
True       0.075110
Name: avg_var_price, dtype: float64
```

```
[73]: print(df_price.groupby('is_weekend')['avg_fix_price'].mean())
```

```
is_weekend
False      20.115581
True       20.136173
Name: avg_fix_price, dtype: float64
```

Overall for the weekends price is not varying much in both energy and power connections.

1.4 Model Training, Hyperparameter Tuning, and Evaluation

In this step we will cover model training, hyperparameter tuning, and evaluation for three types of models: XGBoost Ensemble Method based on Decision Trees, Dense Neural Network based on Dense Layers, and BiDirectional LSTM-CNN model which combines LSTM and CNN layers for Churn Prediction.

While XGBoost works fine with or without Scaled Parameters we will scale the dataset for training as other two models need data scaling. For hyperparameter tuning we will focus on Gradient Boosting and BiLSTM - CNN. For the DNN Model hyperparameter we already tested some combinations in advance and used both regularization with early stopping.

We will evaluate the models on the basis of ROC-AUC Score, F1 Score, Recall, and Accuracy with less focus on Accuracy as ROC-AUC and Recall scores are better predictor for this task.

1.4.1 Data Preperation

In this step we will prepare data for modeling.

```
[74]: #Merging price and client data with left join on price
df_price.drop(columns = 'price_date', inplace = True)
df_merge = df_price.merge(df, on = 'id', how = 'left')
df_merge.head()
```

```
[74]:
```

	id	price_off_peak_var	price_peak_var	\
0	038af19179925da21a25619c5a24b745	0.151367	0.0	
1	038af19179925da21a25619c5a24b745	0.151367	0.0	
2	038af19179925da21a25619c5a24b745	0.151367	0.0	
3	038af19179925da21a25619c5a24b745	0.149626	0.0	
4	038af19179925da21a25619c5a24b745	0.149626	0.0	

	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix	\
0	0.0	44.266931	0.0	0.0	
1	0.0	44.266931	0.0	0.0	
2	0.0	44.266931	0.0	0.0	
3	0.0	44.266931	0.0	0.0	
4	0.0	44.266931	0.0	0.0	

	price_month	is_weekend	is_month_start	...	var_6m_price_mid_peak_var	\
0	1	False	True	...	0.0	
1	2	True	True	...	0.0	
2	3	True	True	...	0.0	
3	4	False	True	...	0.0	
4	5	False	True	...	0.0	

	var_6m_price_off_peak_fix	var_6m_price_peak_fix	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	

3	0.0	0.0
4	0.0	0.0

	var_6m_price_mid_peak_fix	var_6m_price_off_peak	var_6m_price_peak \
0	0.0	0.000003	0.0
1	0.0	0.000003	0.0
2	0.0	0.000003	0.0
3	0.0	0.000003	0.0
4	0.0	0.000003	0.0

	var_6m_price_mid_peak	churn	tenure	end_month
0	0.0	0.0	1096.0	6.0
1	0.0	0.0	1096.0	6.0
2	0.0	0.0	1096.0	6.0
3	0.0	0.0	1096.0	6.0
4	0.0	0.0	1096.0	6.0

[5 rows x 56 columns]

BCGX standard answer did left merge on the client dataset but for this task left merge on price data is an industry norm because it helps us to capture more day to day data representations in the model and predictions.

```
[75]: df_merge.info
```

```
[75]: <bound method DataFrame.info of
price_off_peak_var  price_peak_var  \
0      038af19179925da21a25619c5a24b745      0.151367      0.000000
1      038af19179925da21a25619c5a24b745      0.151367      0.000000
2      038af19179925da21a25619c5a24b745      0.151367      0.000000
3      038af19179925da21a25619c5a24b745      0.149626      0.000000
4      038af19179925da21a25619c5a24b745      0.149626      0.000000
...
192997  16f51cdc2baa19af0b940ee1b3dd17d5      0.119916      0.102232
192998  16f51cdc2baa19af0b940ee1b3dd17d5      0.119916      0.102232
192999  16f51cdc2baa19af0b940ee1b3dd17d5      0.119916      0.102232
193000  16f51cdc2baa19af0b940ee1b3dd17d5      0.119916      0.102232
193001  16f51cdc2baa19af0b940ee1b3dd17d5      0.119916      0.102232

      price_mid_peak_var  price_off_peak_fix  price_peak_fix  \
0      0.000000      44.266931      0.00000
1      0.000000      44.266931      0.00000
2      0.000000      44.266931      0.00000
3      0.000000      44.266931      0.00000
4      0.000000      44.266931      0.00000
...
192997      0.076257      40.728885      24.43733
```

192998	0.076257	40.728885	24.43733
192999	0.076257	40.728885	24.43733
193000	0.076257	40.728885	24.43733
193001	0.076257	40.728885	24.43733

	price_mid_peak_fix	price_month	is_weekend	is_month_start	...	\
0	0.000000	1	False	True	...	
1	0.000000	2	True	True	...	
2	0.000000	3	True	True	...	
3	0.000000	4	False	True	...	
4	0.000000	5	False	True	...	
...	
192997	16.291555	8	True	True	...	
192998	16.291555	9	False	True	...	
192999	16.291555	10	False	True	...	
193000	16.291555	11	True	True	...	
193001	16.291555	12	False	True	...	

	var_6m_price_mid_peak_var	var_6m_price_off_peak_fix	\
0	0.000000e+00	0.0	
1	0.000000e+00	0.0	
2	0.000000e+00	0.0	
3	0.000000e+00	0.0	
4	0.000000e+00	0.0	
...	
192997	4.860000e-10	0.0	
192998	4.860000e-10	0.0	
192999	4.860000e-10	0.0	
193000	4.860000e-10	0.0	
193001	4.860000e-10	0.0	

	var_6m_price_peak_fix	var_6m_price_mid_peak_fix	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	
...	
192997	0.0	0.0	
192998	0.0	0.0	
192999	0.0	0.0	
193000	0.0	0.0	
193001	0.0	0.0	

	var_6m_price_off_peak	var_6m_price_peak	var_6m_price_mid_peak	churn	\
0	0.000003	0.000000	0.000000e+00	0.0	
1	0.000003	0.000000	0.000000e+00	0.0	

2	0.000003	0.000000	0.000000e+00	0.0
3	0.000003	0.000000	0.000000e+00	0.0
4	0.000003	0.000000	0.000000e+00	0.0
...
192997	0.000011	0.000003	4.860000e-10	0.0
192998	0.000011	0.000003	4.860000e-10	0.0
192999	0.000011	0.000003	4.860000e-10	0.0
193000	0.000011	0.000003	4.860000e-10	0.0
193001	0.000011	0.000003	4.860000e-10	0.0

	tenure	end_month
0	1096.0	6.0
1	1096.0	6.0
2	1096.0	6.0
3	1096.0	6.0
4	1096.0	6.0
...
192997	1461.0	6.0
192998	1461.0	6.0
192999	1461.0	6.0
193000	1461.0	6.0
193001	1461.0	6.0

[193002 rows x 56 columns]>

```
[76]: #Checking na values
df_merge.isna().sum()
```

```
[76]: id                                0
price_off_peak_var                      0
price_peak_var                          0
price_mid_peak_var                      0
price_off_peak_fix                      0
price_peak_fix                          0
price_mid_peak_fix                      0
price_month                             0
is_weekend                             0
is_month_start                          0
is_month_end                            0
quarter                                0
var_peak_offpeak_spread                 0
fix_peak_offpeak_spread                 0
avg_var_price                           0
avg_fix_price                           0
peak_var_x_is_weekend                   0
channel_sales                           17853
cons_12m                                17853
```

cons_gas_12m	17853
cons_last_month	17853
forecast_cons_12m	17853
forecast_discount_energy	17853
forecast_meter_rent_12m	17853
forecast_price_energy_off_peak	17853
forecast_price_energy_peak	17853
forecast_price_pow_off_peak	17853
has_gas	17853
imp_cons	17853
margin_gross_pow_ele	17853
nb_prod_act	17853
net_margin	17853
num_years_antig	17853
origin_up	17853
pow_max	17853
var_year_price_off_peak_var	17853
var_year_price_peak_var	17853
var_year_price_mid_peak_var	17853
var_year_price_off_peak_fix	17853
var_year_price_peak_fix	17853
var_year_price_mid_peak_fix	17853
var_year_price_off_peak	17853
var_year_price_peak	17853
var_year_price_mid_peak	17853
var_6m_price_off_peak_var	17853
var_6m_price_peak_var	17853
var_6m_price_mid_peak_var	17853
var_6m_price_off_peak_fix	17853
var_6m_price_peak_fix	17853
var_6m_price_mid_peak_fix	17853
var_6m_price_off_peak	17853
var_6m_price_peak	17853
var_6m_price_mid_peak	17853
churn	17853
tenure	17853
end_month	17853
dtype: int64	

```
[77]: missing_ids = set(df_price['id']) - set(df['id'])
print(f"Number of unmatched client ids: {len(missing_ids)}")
```

Number of unmatched client ids: 1490

```
[78]: #Since Price dataset has more client ids than client df iteself we will only
      ↪cover data for client ids fallinf under sales channels
df_merge = df_merge[df_merge['channel_sales'].notnull()]
```

```
[79]: df_merge.isna().sum()
```

```
[79]: id                                0
      price_off_peak_var                0
      price_peak_var                    0
      price_mid_peak_var                0
      price_off_peak_fix                0
      price_peak_fix                    0
      price_mid_peak_fix                0
      price_month                       0
      is_weekend                        0
      is_month_start                    0
      is_month_end                      0
      quarter                           0
      var_peak_offpeak_spread           0
      fix_peak_offpeak_spread           0
      avg_var_price                     0
      avg_fix_price                     0
      peak_var_x_is_weekend             0
      channel_sales                     0
      cons_12m                          0
      cons_gas_12m                      0
      cons_last_month                   0
      forecast_cons_12m                  0
      forecast_discount_energy           0
      forecast_meter_rent_12m            0
      forecast_price_energy_off_peak     0
      forecast_price_energy_peak         0
      forecast_price_pow_off_peak        0
      has_gas                           0
      imp_cons                           0
      margin_gross_pow_ele               0
      nb_prod_act                        0
      net_margin                         0
      num_years_antig                    0
      origin_up                          0
      pow_max                            0
      var_year_price_off_peak_var        0
      var_year_price_peak_var            0
      var_year_price_mid_peak_var        0
      var_year_price_off_peak_fix        0
      var_year_price_peak_fix            0
      var_year_price_mid_peak_fix        0
      var_year_price_off_peak            0
      var_year_price_peak                0
      var_year_price_mid_peak            0
      var_6m_price_off_peak_var          0
```

```

var_6m_price_peak_var          0
var_6m_price_mid_peak_var      0
var_6m_price_off_peak_fix      0
var_6m_price_peak_fix          0
var_6m_price_mid_peak_fix      0
var_6m_price_off_peak          0
var_6m_price_peak              0
var_6m_price_mid_peak          0
churn                          0
tenure                         0
end_month                      0
dtype: int64

```

```
[80]: df_merge.select_dtypes(include = ['object', 'category']).head()
```

```

[80]:
           id          channel_sales \
0  038af19179925da21a25619c5a24b745  foosdfpfkusacimwkcsosbicdxkicaua
1  038af19179925da21a25619c5a24b745  foosdfpfkusacimwkcsosbicdxkicaua
2  038af19179925da21a25619c5a24b745  foosdfpfkusacimwkcsosbicdxkicaua
3  038af19179925da21a25619c5a24b745  foosdfpfkusacimwkcsosbicdxkicaua
4  038af19179925da21a25619c5a24b745  foosdfpfkusacimwkcsosbicdxkicaua

           origin_up
0  ldkssxwpmemidmecebumciepifcamkci
1  ldkssxwpmemidmecebumciepifcamkci
2  ldkssxwpmemidmecebumciepifcamkci
3  ldkssxwpmemidmecebumciepifcamkci
4  ldkssxwpmemidmecebumciepifcamkci

```

Now we will remove id variables. For channel_sales and origin_up we will use one hot encoding.

```
[81]: df_merge = pd.get_dummies(df_merge, columns=['channel_sales', 'origin_up'],
    ↪drop_first=True, dtype=int)
```

One important note, in case of imbalance data generally oversampling on minority class is done using SMOTE however Gradient Boosting techniques are better suited for such dataset and generally SMOTE is not considered useful overall.

```

[87]: #In order to avoid possible data leakage we will split the dataset on the basis
    ↪of clients with 20% clients going into test instead of 20% rows randomly
unique_customers = df_merge['id'].unique()
train_customers, test_customers = train_test_split(unique_customers, test_size=
    ↪0.2, random_state = 10)
train_df = df_merge[df_merge['id'].isin(train_customers)]
test_df = df_merge[df_merge['id'].isin(test_customers)]

```

```
[93]: train_df.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 140127 entries, 12 to 193001

Data columns (total 65 columns):

#	Column	Non-Null Count	Dtype
0	price_off_peak_var	140127 non-null	float64
1	price_peak_var	140127 non-null	float64
2	price_mid_peak_var	140127 non-null	float64
3	price_off_peak_fix	140127 non-null	float64
4	price_peak_fix	140127 non-null	float64
5	price_mid_peak_fix	140127 non-null	float64
6	price_month	140127 non-null	int32
7	is_weekend	140127 non-null	bool
8	is_month_start	140127 non-null	bool
9	is_month_end	140127 non-null	bool
10	quarter	140127 non-null	int32
11	var_peak_offpeak_spread	140127 non-null	float64
12	fix_peak_offpeak_spread	140127 non-null	float64
13	avg_var_price	140127 non-null	float64
14	avg_fix_price	140127 non-null	float64
15	peak_var_x_is_weekend	140127 non-null	float64
16	cons_12m	140127 non-null	float64
17	cons_gas_12m	140127 non-null	float64
18	cons_last_month	140127 non-null	float64
19	forecast_cons_12m	140127 non-null	float64
20	forecast_discount_energy	140127 non-null	float64
21	forecast_meter_rent_12m	140127 non-null	float64
22	forecast_price_energy_off_peak	140127 non-null	float64
23	forecast_price_energy_peak	140127 non-null	float64
24	forecast_price_pow_off_peak	140127 non-null	float64
25	has_gas	140127 non-null	float64
26	imp_cons	140127 non-null	float64
27	margin_gross_pow_ele	140127 non-null	float64
28	nb_prod_act	140127 non-null	float64
29	net_margin	140127 non-null	float64
30	num_years_antig	140127 non-null	float64
31	pow_max	140127 non-null	float64
32	var_year_price_off_peak_var	140127 non-null	float64
33	var_year_price_peak_var	140127 non-null	float64
34	var_year_price_mid_peak_var	140127 non-null	float64
35	var_year_price_off_peak_fix	140127 non-null	float64
36	var_year_price_peak_fix	140127 non-null	float64
37	var_year_price_mid_peak_fix	140127 non-null	float64
38	var_year_price_off_peak	140127 non-null	float64
39	var_year_price_peak	140127 non-null	float64
40	var_year_price_mid_peak	140127 non-null	float64
41	var_6m_price_off_peak_var	140127 non-null	float64
42	var_6m_price_peak_var	140127 non-null	float64

```

43 var_6m_price_mid_peak_var          140127 non-null float64
44 var_6m_price_off_peak_fix          140127 non-null float64
45 var_6m_price_peak_fix              140127 non-null float64
46 var_6m_price_mid_peak_fix          140127 non-null float64
47 var_6m_price_off_peak              140127 non-null float64
48 var_6m_price_peak                  140127 non-null float64
49 var_6m_price_mid_peak              140127 non-null float64
50 churn                              140127 non-null float64
51 tenure                             140127 non-null float64
52 end_month                           140127 non-null float64
53 channel_sales_epumfxlbckeskwexbiuasklxalciuu 140127 non-null int64
54 channel_sales_ewpakwlliwisiwduibdlfmalxowmwpci 140127 non-null int64
55 channel_sales_fixdbufsefwooaasfcxdxadsiekocaaa 140127 non-null int64
56 channel_sales_foosdfpfkusacimwkcsofbicdxkicaua 140127 non-null int64
57 channel_sales_lmkebamcaaclubfxadlmueccxoimlema 140127 non-null int64
58 channel_sales_sddiedcslfslkckwlfkdpoeailfpeds 140127 non-null int64
59 channel_sales_usilxuppasemubllpokaafesmlibmsdf 140127 non-null int64
60 origin_up_ewxeelcelemmiwuafmddpobolfuxioce    140127 non-null int64
61 origin_up_kamkkxfxxuwbdslkwifmmcsiusiuosws    140127 non-null int64
62 origin_up_ldkssxwpmemidmecebumciepifcamkci    140127 non-null int64
63 origin_up_lxidpiddsbxsbosboudacockeimpuepw     140127 non-null int64
64 origin_up_usapbecpfoloekilkwsdiboslwaxobdp     140127 non-null int64
dtypes: bool(3), float64(48), int32(2), int64(12)
memory usage: 66.7 MB

```

```

[339]: X_train = train_df.drop(columns = 'churn')
        Y_train = train_df['churn']
        X_test = test_df.drop(columns = 'churn')
        Y_test = test_df['churn']
        print(X_train.shape)
        print(Y_train.shape)
        print(X_test.shape)
        print(Y_test.shape)

```

```

(140127, 64)
(140127,)
(35022, 64)
(35022,)

```

```

[340]: scaler = StandardScaler() #normalizing skewed variables and other numerical
        ↪variables
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)

```

```

[155]: model_xg = XGBClassifier(n_estimators = 100,
                                learning_rate = 0.1,
                                max_depth = 8,

```

```

        random_state = 10,
        subsample = 0.8,
        colsample_bytree = 0.8,
        scale_pos_weight = 9)

```

```
[162]: model_xg.fit(X_train_scaled, Y_train)
```

```
[162]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bynode=None,
                    colsample_bytree=0.8, device=None, early_stopping_rounds=None,
                    enable_categorical=False, eval_metric=None, feature_types=None,
                    feature_weights=None, gamma=None, grow_policy=None,
                    importance_type=None, interaction_constraints=None,
                    learning_rate=0.1, max_bin=None, max_cat_threshold=None,
                    max_cat_to_onehot=None, max_delta_step=None, max_depth=8,
                    max_leaves=None, min_child_weight=None, missing=nan,
                    monotone_constraints=None, multi_strategy=None, n_estimators=100,
                    n_jobs=None, num_parallel_tree=None, ...)

```


```
[163]: y_pred = model_xg.predict(X_test_scaled)
y_pred_proba = model_xg.predict_proba(X_test_scaled)[:, 1]

fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
roc_auc = roc_auc_score(Y_test, y_pred_proba)

# Plot ROC curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for XGBoost Churn Model')
plt.legend(loc='lower right')
plt.show()

print(f"ROC AUC Score: {roc_auc:.2f}")

```



ROC AUC Score: 0.68

```
[164]: # Classification report
print(classification_report(Y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(Y_test, y_pred)
print("Confusion Matrix:\n", cm)
print(f'F1 Score: {f1_score(Y_test, y_pred)}')
```

	precision	recall	f1-score	support
0.0	0.92	0.94	0.93	31649
1.0	0.32	0.28	0.30	3373
accuracy			0.87	35022
macro avg	0.62	0.61	0.61	35022
weighted avg	0.87	0.87	0.87	35022

Confusion Matrix:


```
[[29682 1967]
 [ 2438  935]]
F1 Score: 0.29800796812749003
```

Performing hyperparameter tuning for XGboost model.

```
[171]: param_grid = {"n_estimators": [100, 200, 300], #number of helper models
                    "learning_rate": [0.1, 0.05, 0.01], #contribution rate for helper
                    ↪models
                    "max_depth": [3, 5, 8], #max depth of tree
                    "subsample": [0.7, 0.8, 1.0], #percentage of training data
                    ↪available
                    "colsample_bytree": [0.7, 0.8, 1.0], #percentage of features
                    ↪available
                    "scale_pos_weight": [9] #used for unbalanced dataset
                }

grid_search_xg = GridSearchCV(
    estimator = XGBClassifier(eval_metric = 'logloss'),
    param_grid = param_grid,
    scoring = 'roc_auc',
    cv = 5,
    verbose = 1,
    n_jobs = -1
)

grid_search_xg.fit(X_train_scaled, Y_train)
```

Fitting 5 folds for each of 243 candidates, totalling 1215 fits

```
[171]: GridSearchCV(cv=5,
                    estimator=XGBClassifier(base_score=None, booster=None,
                                             callbacks=None, colsample_bylevel=None,
                                             colsample_bynode=None,
                                             colsample_bytree=None, device=None,
                                             early_stopping_rounds=None,
                                             enable_categorical=False,
                                             eval_metric='logloss', feature_types=None,
                                             feature_weights=None, gamma=None,
                                             grow_policy=None, importance_type=None,
                                             interaction_constraint...
                                             max_leaves=None, min_child_weight=None,
                                             missing=nan, monotone_constraints=None,
                                             multi_strategy=None, n_estimators=None,
                                             n_jobs=None, num_parallel_tree=None, ...),
                    n_jobs=-1,
                    param_grid={'colsample_bytree': [0.7, 0.8, 1.0],
                                'learning_rate': [0.1, 0.05, 0.01],
```

```

        'max_depth': [3, 5, 8],
        'n_estimators': [100, 200, 300],
        'scale_pos_weight': [9],
        'subsample': [0.7, 0.8, 1.0]},
    scoring='roc_auc', verbose=1)

```

```

[176]: print(grid_search_xg.best_params_)
       print(grid_search_xg.best_score_)

```

```

{'colsample_bytree': 0.7, 'learning_rate': 0.01, 'max_depth': 5, 'n_estimators':
300, 'scale_pos_weight': 9, 'subsample': 0.7}
0.6906782885031666

```

```

[173]: best_model_xg = grid_search_xg.best_estimator_

```


```

[177]: y_pred_proba = best_model_xg.predict_proba(X_test_scaled)[: ,1]
       fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
       roc_auc = roc_auc_score(Y_test, y_pred_proba)

       # Plot ROC curve
       plt.figure(figsize=(8,6))
       plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
       plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')
       plt.xlabel('False Positive Rate')
       plt.ylabel('True Positive Rate')
       plt.title('ROC Curve for XGBoost Churn Model')
       plt.legend(loc='lower right')
       plt.show()

       print(f"ROC AUC Score: {roc_auc:.2f}")

```



ROC AUC Score: 0.69

[]:

```
[178]: # Predict class labels (default threshold 0.5)
y_pred = best_model_xg.predict(X_test_scaled)

# Classification report
print(classification_report(Y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(Y_test, y_pred)
print("Confusion Matrix:\n", cm)
print(f'F1 Score: {f1_score(Y_test, y_pred)}')
```

	precision	recall	f1-score	support
0.0	0.94	0.76	0.84	31649
1.0	0.19	0.51	0.27	3373

accuracy			0.74	35022
macro avg	0.56	0.64	0.56	35022
weighted avg	0.86	0.74	0.79	35022

Confusion Matrix:

```
[[24195  7454]
 [ 1649  1724]]
```

F1 Score: 0.2747191458847901

State of the art industry Churn prediction models have f1 score of 0.5 and more with roc-auc score above 0.8. While our F1 score may look low, Churn prediction is a highly complex task in real world complex datasets and most companies start with models with such an F1 score and improve it further. In our case, our model is already far better than what BCGX proposed for this task. Their model was based on decision trees and had f1 score of 0.09 and recall of 0.05.


In contrast to that model our model is leagues ahead with 0.27 F1 and 0.51 recall. Our hyperparameter tuning actually improved in recall as we focused on roc-auc score. If we focused on F1 Score, we would have gotten 0.28 F1 score and 0.33 recall. Generally, both F1 score and ROC-AUC score should be improved further by feature engineering and adding more data but we can't do them in our current project as we are limited to dataset obtained from BCGX.

```
[185]: feature_importance = pd.DataFrame({"features": X_train.columns, 'importance':
    ↪best_model_xg.feature_importances_})
feature_importance.sort_values('importance', ascending = False).head()
```

```
[185]:
```

	features	importance
60	origin_up_kamkkxfxxuwbdslkwifmmcsiusuosws	0.052923
62	origin_up_lxidpiddsbxsbosboudacockeimpuepw	0.050283
27	margin_gross_pow_ele	0.043019
51	end_month	0.032805
14	avg_fix_price	0.028505

```
[187]: plt.figure(figsize=(15, 25))
plt.title('Feature Importances')
plt.barh(range(len(feature_importance)), feature_importance['importance'],
    ↪color='b', align='center')
plt.yticks(range(len(feature_importance)), feature_importance['features'])
plt.xlabel('Importance')
plt.show()
```



Based on the feature importance we can observe that Origin_up grouping is a major factor followed by gross margin, end month, and average price.

So, is price a good indicator for the churn prediction for PowerCO? On the basis of the above analysis, for our model top contributing factor is origin_up group and gross margin. Average price,

one of the price indicator, has relative importance of 2.8% so its a relatively moderate contributor, not a top contributing factor. However, we need better feature engineer and more data to be sure.

A lot of predictors have near 0 feature importance so we would remove them for our final Xgboost model. However, while the model needs more feature engineering and data, after our analysis we can confirm that this dataset itself is limited and not usable for more feature engineering.

```
[213]: X_train_xgb = X_train.drop(columns = list(feature_importance.
    ↪loc[feature_importance['importance'] == 0,].iloc[:,0]))
X_test_xgb = X_test.drop(columns = list(feature_importance.
    ↪loc[feature_importance['importance'] == 0,].iloc[:,0]))

#We will not scale the parameters this time as XG Boost is invariant to it_
    ↪overall, we scaled to fit both XGboost and neurel network
```

```
[218]: model_xgb = XGBClassifier(eval_metric = 'logloss',
    colsample_bytree = 0.7,
    learning_rate = 0.01,
    max_depth = 5,
    n_estimators = 300,
    scale_pos_weight = 9,
    subsample = 0.7)
model_xgb.fit(X_train_xgb, Y_train)
```

```
[218]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=0.7, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric='logloss',
    feature_types=None, feature_weights=None, gamma=None,
    grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.01, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=5, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=300, n_jobs=None,
    num_parallel_tree=None, ...)
```


```
[219]: y_pred = model_xgb.predict(X_test_xgb)
y_pred_proba = model_xgb.predict_proba(X_test_xgb)[: , 1]

fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
roc_auc = roc_auc_score(Y_test, y_pred_proba)

# Plot ROC curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for XGBoost Churn Model')
plt.legend(loc='lower right')
plt.show()
```

```
print(f"ROC AUC Score: {roc_auc:.2f}")
```



ROC AUC Score: 0.68

```
[220]: # Predict class labels (default threshold 0.5)
y_pred = model_xgb.predict(X_test_xgb)

# Classification report
print(classification_report(Y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(Y_test, y_pred)
print("Confusion Matrix:\n", cm)
print(f'F1 Score: {f1_score(Y_test, y_pred)}')
```

	precision	recall	f1-score	support
0.0	0.93	0.76	0.84	31649
1.0	0.18	0.50	0.27	3373
accuracy			0.74	35022
macro avg	0.56	0.63	0.55	35022
weighted avg	0.86	0.74	0.79	35022

Confusion Matrix:

```
[[24182  7467]
 [ 1692 1681]]
```


F1 Score: 0.26850890503953356

```
[222]: feature_importance = pd.DataFrame({"features": X_train_xgb.columns,
    ↳ 'importance': model_xgb.feature_importances_})
feature_importance.sort_values('importance', ascending = False).head()
```

```
[222]:
```

	features	importance
52	origin_up_lxidpiddsbxsbosboudacockeimpuepw	0.058003
50	origin_up_kamkkxfxxuwbdslkwifmmcsiusiosws	0.044520
21	margin_gross_pow_ele	0.041389
9	avg_fix_price	0.033047
45	end_month	0.032014

```
[223]: plt.figure(figsize=(15, 25))
plt.title('Feature Importances')
plt.barh(range(len(feature_importance)), feature_importance['importance'],
    ↳ color='b', align='center')
plt.yticks(range(len(feature_importance)), feature_importance['features'])
plt.xlabel('Importance')
plt.show()
```

Final XGBoost model has slightly lower performance but its not that different overall as its overall ROC drops by 1% to 68%. We can use either models as both of them are better than BCGX model answer where we had 52% ROC-AUC Score.

For feature importance we can observed that by removing redundant variables importance for most

important variable did improve. Overall, PowerCo was concerned with Churn Analysis and if price is a major churn variable. In BCGX model answer price predictors had around 2% importance but in ours we have price variable in top 5 most important feature.

Still, we can't say definitely that price impact churn in an impactful way like origin_up classification and gross margin, but still it's a moderately impactful predictor for churn which does warrant further analysis.

1.4.2 Neurel Networks

In this step we will fit two tensorflow powered deep learning models. One will be based on Dense Layers based on relu activation while other will be based on BiDirectional LSTM-CNN Model.

For DNN model we will use relu activations throughout the model with proper batch normalization and weights initializer to avoid possible gradient exploding/shrinkage. And for BiLSTM-CNN model we will be using padding and masking for the time series data.

```
[228]: #Checking GPU availability

print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))

hello = tf.constant("Hello TensorFlow!")
print(hello.numpy())
```

```
Num GPUs Available:  1
b'Hello TensorFlow!'
```

```
[232]: X_train.shape[1]
```

```
[232]: 64
```

```
[248]: #DNN Based model

model_dnn = Sequential([
    Dense(32, input_dim = X_train.shape[1], kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
    Dense(64, kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
    Dense(128, kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
    Dense(64, kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
    Dense(32, kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
```

```

    Dense(16, kernel_initializer = HeNormal()),
    BatchNormalization(),
    Activation('relu'),
    Dense(1, activation = 'linear')
])

```

```
[249]: model_dnn.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 32)	2,080
batch_normalization_12 (BatchNormalization)	(None, 32)	128
activation_12 (Activation)	(None, 32)	0
dense_16 (Dense)	(None, 64)	2,112
batch_normalization_13 (BatchNormalization)	(None, 64)	256
activation_13 (Activation)	(None, 64)	0
dense_17 (Dense)	(None, 128)	8,320
batch_normalization_14 (BatchNormalization)	(None, 128)	512
activation_14 (Activation)	(None, 128)	0
dense_18 (Dense)	(None, 64)	8,256
batch_normalization_15 (BatchNormalization)	(None, 64)	256
activation_15 (Activation)	(None, 64)	0
dense_19 (Dense)	(None, 32)	2,080
batch_normalization_16 (BatchNormalization)	(None, 32)	128
activation_16 (Activation)	(None, 32)	0

dense_20 (Dense)	(None, 16)	528
batch_normalization_17 (BatchNormalization)	(None, 16)	64
activation_17 (Activation)	(None, 16)	0
dense_21 (Dense)	(None, 1)	17

Total params: 24,737 (96.63 KB)

Trainable params: 24,065 (94.00 KB)

Non-trainable params: 672 (2.62 KB)

```
[247]: early_stopping = EarlyStopping(monitor = 'val_AUC', patience = 10,
    ↪restore_best_weights=True)
```

```
[250]: model_dnn.compile(optimizer = tf.keras.optimizers.Adam(learning_rate = 0.001),
    loss = tf.keras.losses.BinaryCrossentropy(from_logits = True),
    metrics = ['accuracy', 'AUC'])
model_dnn.fit(X_train_scaled, Y_train, epochs = 100, batch_size = 32,
    ↪validation_split = 0.2, callbacks = [early_stopping])
```

```
Epoch 1/100
3504/3504          23s 6ms/step -
AUC: 0.5030 - accuracy: 0.8685 - loss: 0.4226 - val_AUC: 0.5008 - val_accuracy:
0.9114 - val_loss: 0.2919
Epoch 2/100
3504/3504          17s 5ms/step -
AUC: 0.5277 - accuracy: 0.9022 - loss: 0.2915 - val_AUC: 0.5298 - val_accuracy:
0.9094 - val_loss: 0.2960
Epoch 3/100
3504/3504          17s 5ms/step -
AUC: 0.5614 - accuracy: 0.9070 - loss: 0.2739 - val_AUC: 0.5341 - val_accuracy:
0.9092 - val_loss: 0.2963
Epoch 4/100
3504/3504          16s 5ms/step -
AUC: 0.5848 - accuracy: 0.9120 - loss: 0.2571 - val_AUC: 0.5480 - val_accuracy:
0.9114 - val_loss: 0.3088
Epoch 5/100
3504/3504          17s 5ms/step -
AUC: 0.6026 - accuracy: 0.9128 - loss: 0.2520 - val_AUC: 0.5701 - val_accuracy:
```

0.9137 - val_loss: 0.3099
 Epoch 6/100
 3504/3504 16s 5ms/step -
 AUC: 0.6247 - accuracy: 0.9183 - loss: 0.2343 - val_AUC: 0.5574 - val_accuracy:
 0.9111 - val_loss: 0.3295
 Epoch 7/100
 3504/3504 16s 5ms/step -
 AUC: 0.6366 - accuracy: 0.9203 - loss: 0.2264 - val_AUC: 0.5489 - val_accuracy:
 0.9062 - val_loss: 0.3263
 Epoch 8/100
 3504/3504 16s 5ms/step -
 AUC: 0.6508 - accuracy: 0.9225 - loss: 0.2197 - val_AUC: 0.5708 - val_accuracy:
 0.9128 - val_loss: 0.3156
 Epoch 9/100
 3504/3504 16s 5ms/step -
 AUC: 0.6600 - accuracy: 0.9244 - loss: 0.2155 - val_AUC: 0.5588 - val_accuracy:
 0.9138 - val_loss: 0.3434
 Epoch 10/100
 3504/3504 16s 5ms/step -
 AUC: 0.6644 - accuracy: 0.9246 - loss: 0.2111 - val_AUC: 0.5746 - val_accuracy:
 0.9120 - val_loss: 0.3391
 Epoch 11/100
 3504/3504 18s 5ms/step -
 AUC: 0.6746 - accuracy: 0.9279 - loss: 0.2043 - val_AUC: 0.5753 - val_accuracy:
 0.9137 - val_loss: 0.3377
 Epoch 12/100
 3504/3504 17s 5ms/step -
 AUC: 0.6804 - accuracy: 0.9279 - loss: 0.2022 - val_AUC: 0.5797 - val_accuracy:
 0.9170 - val_loss: 0.3515
 Epoch 13/100
 3504/3504 17s 5ms/step -
 AUC: 0.6859 - accuracy: 0.9276 - loss: 0.2000 - val_AUC: 0.5686 - val_accuracy:
 0.9094 - val_loss: 0.3631
 Epoch 14/100
 3504/3504 18s 5ms/step -
 AUC: 0.6944 - accuracy: 0.9307 - loss: 0.1913 - val_AUC: 0.5771 - val_accuracy:
 0.9031 - val_loss: 0.3715
 Epoch 15/100
 3504/3504 17s 5ms/step -
 AUC: 0.7000 - accuracy: 0.9307 - loss: 0.1898 - val_AUC: 0.5722 - val_accuracy:
 0.9102 - val_loss: 0.3787
 Epoch 16/100
 3504/3504 17s 5ms/step -
 AUC: 0.6993 - accuracy: 0.9305 - loss: 0.1916 - val_AUC: 0.5836 - val_accuracy:
 0.9115 - val_loss: 0.3704
 Epoch 17/100
 3504/3504 17s 5ms/step -
 AUC: 0.7095 - accuracy: 0.9325 - loss: 0.1820 - val_AUC: 0.5735 - val_accuracy:

```

0.9081 - val_loss: 0.3750
Epoch 18/100
3504/3504          17s 5ms/step -
AUC: 0.7135 - accuracy: 0.9326 - loss: 0.1812 - val_AUC: 0.5781 - val_accuracy:
0.9093 - val_loss: 0.3796
Epoch 19/100
3504/3504          17s 5ms/step -
AUC: 0.7154 - accuracy: 0.9326 - loss: 0.1815 - val_AUC: 0.5689 - val_accuracy:
0.9110 - val_loss: 0.4018
Epoch 20/100
3504/3504          18s 5ms/step -
AUC: 0.7248 - accuracy: 0.9348 - loss: 0.1753 - val_AUC: 0.5722 - val_accuracy:
0.9069 - val_loss: 0.3962
Epoch 21/100
3504/3504          17s 5ms/step -
AUC: 0.7266 - accuracy: 0.9356 - loss: 0.1741 - val_AUC: 0.5645 - val_accuracy:
0.9075 - val_loss: 0.4293
Epoch 22/100
3504/3504          17s 5ms/step -
AUC: 0.7317 - accuracy: 0.9354 - loss: 0.1721 - val_AUC: 0.5714 - val_accuracy:
0.9105 - val_loss: 0.3909
Epoch 23/100
3504/3504          17s 5ms/step -
AUC: 0.7321 - accuracy: 0.9361 - loss: 0.1712 - val_AUC: 0.5771 - val_accuracy:
0.9096 - val_loss: 0.4071
Epoch 24/100
3504/3504          17s 5ms/step -
AUC: 0.7361 - accuracy: 0.9368 - loss: 0.1697 - val_AUC: 0.5543 - val_accuracy:
0.9029 - val_loss: 0.4323
Epoch 25/100
3504/3504          17s 5ms/step -
AUC: 0.7387 - accuracy: 0.9376 - loss: 0.1673 - val_AUC: 0.5680 - val_accuracy:
0.9028 - val_loss: 0.4226
Epoch 26/100
3504/3504          17s 5ms/step -
AUC: 0.7472 - accuracy: 0.9395 - loss: 0.1616 - val_AUC: 0.5650 - val_accuracy:
0.9104 - val_loss: 0.4144

```

[250]: <keras.src.callbacks.history.History at 0x7695c04c7dd0>

```

[ ]: y_pred = model_dnn.predict(X_test_xgb)
     y_pred_proba = model_xgb.predict_proba(X_test_xgb)[: , 1]

     fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
     roc_auc = roc_auc_score(Y_test, y_pred_proba)

     # Plot ROC curve

```

```

plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for XGBoost Churn Model')
plt.legend(loc='lower right')
plt.show()

print(f"ROC AUC Score: {roc_auc:.2f}")

```

```

[341]: y_pred = tf.nn.sigmoid(model_dnn.predict(X_test_scaled)) > 0.5
y_pred_proba = tf.nn.sigmoid(model_dnn.predict(X_test_scaled))

fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
roc_auc = roc_auc_score(Y_test, y_pred_proba)

# Plot ROC curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for XGBoost Churn Model')
plt.legend(loc='lower right')
plt.show()


print(f"ROC AUC Score: {roc_auc:.2f}")

```

```

1095/1095          2s 2ms/step
1095/1095          2s 1ms/step

```



ROC AUC Score: 0.62

```
[264]: # Classification report
print(classification_report(Y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(Y_test, y_pred)
print("Confusion Matrix:\n", cm)
print(f'F1 Score: {f1_score(Y_test, y_pred)}')
```

	precision	recall	f1-score	support
0.0	0.92	0.97	0.94	31649
1.0	0.38	0.19	0.25	3373
accuracy			0.89	35022
macro avg	0.65	0.58	0.60	35022
weighted avg	0.87	0.89	0.88	35022

Confusion Matrix:


```
[[30632  1017]
 [ 2742   631]]
F1 Score: 0.2513443537143995
```

Our DNN model didn't work better than gradient boosting. Even when focus was done on AUC score on training we got a model with very less recall overall compared to XGBoosting algorithm.

Now we will fit a BiLSTM-CNN model based on the research paper: Customer churn prediction using composite deep learning technique Link: <https://www.nature.com/articles/s41598-023-44396-w#:~:text=A%20composite%20deep%20learning%20model,dataset%20for%20predicting%20customer%20attrition>

For this model we need timestamps as its based on LSTM layers and since we have 12 month data for each clients we can use this model effectively. Only concern here is the lack of proper 12 month data for churned clients as we can not source more data from BCGX, to fill this gap we will use Zero Padding for churned clients missing some monthly data and proper masking during the model training.

```
[300]: # Ensure data is sorted by client and month
df_merge_dm = df_merge.sort_values(['id', 'price_month'])

# List of all clients
clients = df_merge_dm['id'].unique()
num_clients = len(clients)
num_months = 12
num_features = 63

# List of feature columns (replace with your actual feature names)
feature_cols = [col for col in df_merge_dm.columns if col not in ['id', 'price_month', 'churn']]

# Initialize arrays
X = np.zeros((num_clients, num_months, num_features))
y = np.zeros(num_clients)

for i, client in enumerate(clients):
    client_data = df_merge_dm[df_merge_dm['id'] == client].
    ↪sort_values('price_month')
    # Get feature values
    features = client_data[feature_cols].values
    # Zero pad if less than 12 months
    if features.shape[0] < num_months:
        padded = np.zeros((num_months, num_features))
        padded[-features.shape[0]:, :] = features # pad at the start, keep
    ↪recent months at the end
    else:
        padded = features[-num_months:, :] # take last 12 months if more
    X[i] = padded
    y[i] = client_data['churn'].iloc[-1]
```

```
[298]: #BiLSTM-CNN Model. It is based on the research paper mentioned before

inputs = Input(shape=(num_months, num_features))
x = Masking(mask_value=0.0)(inputs)
x = Bidirectional(LSTM(64, return_sequences = True))(x)
x = Conv1D(filters = 32, kernel_size = 3, activation = 'relu', padding = '
↳same')(x)
x = GlobalMaxPooling1D()(x)
x = Dense(64, activation = 'relu')(x)
x = Dropout(0.3)(x)
outputs = Dense(1, activation = 'sigmoid')(x)

model = Model(inputs, outputs)
model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = '
↳['accuracy', 'AUC'])
model.summary()
```

```
[298]: (175149, 66)
```

```
[308]: X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size = 0.2,
↳random_state = 10)
```

```
[316]: num_clients, num_months, num_features = X_train.shape

# Reshape to 2D for scaling
X_2d = X_train.reshape(-1, num_features) # shape: (num_clients * num_months,
↳num_features)

# Fit scaler on training data only!
scaler = StandardScaler()
X_2d_scaled = scaler.fit_transform(X_2d)

# Reshape back to 3D
X_train_scaled = X_2d_scaled.reshape(num_clients, num_months, num_features)
```

```
[329]: timesteps = 12
n_features = 63

inputs = tf.keras.Input(shape=(timesteps, n_features))

inputs = Input(shape=(num_months, num_features))
x = Masking(mask_value=0.0)(inputs) #masking for padded data
x = Bidirectional(LSTM(64, return_sequences=True))(x) #Bi-directional LSTM
↳layer
x = Conv1D(filters=32, kernel_size=3, activation='relu', padding='same')(x)
↳#Convolutional Layer
x = GlobalMaxPooling1D()(x)
```

```

x = Dense(64, activation='relu')(x)
x = Dropout(0.3)(x)
outputs = Dense(1, activation='sigmoid')(x)

model = models.Model(inputs, outputs)
model.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate = 0.001),
    loss = 'binary_crossentropy',
    metrics=['accuracy', tf.keras.metrics.AUC(name='auc'), tf.keras.metrics.
↳Precision(), tf.keras.metrics.Recall()]
)
model.summary()

```

Model: "functional_7"

Layer (type)	Output Shape	Param #	Connected to
input_layer_12 (InputLayer)	(None, 12, 63)	0	-
not_equal_2 (NotEqual)	(None, 12, 63)	0	input_layer_12[0...
masking_2 (Masking)	(None, 12, 63)	0	input_layer_12[0...
any_2 (Any)	(None, 12)	0	not_equal_2[0][0]
bidirectional_5 (Bidirectional)	(None, 12, 128)	65,536	masking_2[0][0], any_2[0][0]
conv1d_5 (Conv1D)	(None, 12, 32)	12,320	bidirectional_5[...
global_max_pooling... (GlobalMaxPooling1...)	(None, 32)	0	conv1d_5[0][0]
dense_32 (Dense)	(None, 64)	2,112	global_max_pooli...
dropout_5 (Dropout)	(None, 64)	0	dense_32[0][0]
dense_33 (Dense)	(None, 1)	65	dropout_5[0][0]

Total params: 80,033 (312.63 KB)

Trainable params: 80,033 (312.63 KB)

Non-trainable params: 0 (0.00 B)

```
[330]: history = model.fit(X_train_scaled, Y_train, epochs = 50,  
                           batch_size = 32, validation_split = 0.2,  
                           callbacks = [EarlyStopping(monitor = 'val_AUC', patience = 10,  
↪10, restore_best_weights = True, mode = 'max'))]
```

Epoch 1/50

293/293 10s 30ms/step -

accuracy: 0.8871 - auc: 0.5489 - loss: 0.3576 - precision_5: 0.0792 - recall_5:
0.0198 - val_accuracy: 0.9050 - val_auc: 0.6017 - val_loss: 0.3111 -
val_precision_5: 0.0000e+00 - val_recall_5: 0.0000e+00

Epoch 2/50

293/293 8s 26ms/step -

accuracy: 0.9043 - auc: 0.6372 - loss: 0.3060 - precision_5: 0.0000e+00 -
recall_5: 0.0000e+00 - val_accuracy: 0.9050 - val_auc: 0.6376 - val_loss: 0.3082
- val_precision_5: 0.0000e+00 - val_recall_5: 0.0000e+00

Epoch 3/50

293/293 8s 26ms/step -

accuracy: 0.9076 - auc: 0.6641 - loss: 0.2959 - precision_5: 0.0000e+00 -
recall_5: 0.0000e+00 - val_accuracy: 0.9050 - val_auc: 0.6300 - val_loss: 0.3084
- val_precision_5: 0.0000e+00 - val_recall_5: 0.0000e+00

Epoch 4/50

293/293 8s 27ms/step -

accuracy: 0.9074 - auc: 0.7159 - loss: 0.2814 - precision_5: 0.6334 - recall_5:
0.0051 - val_accuracy: 0.9029 - val_auc: 0.6440 - val_loss: 0.3053 -
val_precision_5: 0.2222 - val_recall_5: 0.0090

Epoch 5/50

293/293 8s 27ms/step -

accuracy: 0.9046 - auc: 0.7129 - loss: 0.2895 - precision_5: 0.5843 - recall_5:
0.0476 - val_accuracy: 0.9046 - val_auc: 0.6297 - val_loss: 0.3113 -
val_precision_5: 0.4545 - val_recall_5: 0.0225

Epoch 6/50

293/293 8s 27ms/step -

accuracy: 0.9073 - auc: 0.7175 - loss: 0.2846 - precision_5: 0.6549 - recall_5:
0.0877 - val_accuracy: 0.9042 - val_auc: 0.6307 - val_loss: 0.3186 -
val_precision_5: 0.4688 - val_recall_5: 0.0676

Epoch 7/50

293/293 7s 25ms/step -

accuracy: 0.9106 - auc: 0.7367 - loss: 0.2751 - precision_5: 0.7484 - recall_5:
0.1178 - val_accuracy: 0.9059 - val_auc: 0.6182 - val_loss: 0.3173 -
val_precision_5: 0.5333 - val_recall_5: 0.0721

Epoch 8/50

293/293 8s 27ms/step -

accuracy: 0.9132 - auc: 0.7715 - loss: 0.2635 - precision_5: 0.7843 - recall_5:

0.1779 - val_accuracy: 0.9054 - val_auc: 0.6553 - val_loss: 0.3077 -
val_precision_5: 0.5152 - val_recall_5: 0.0766
Epoch 9/50
293/293 8s 26ms/step -
accuracy: 0.9156 - auc: 0.7750 - loss: 0.2559 - precision_5: 0.7717 - recall_5:
0.1370 - val_accuracy: 0.9024 - val_auc: 0.6663 - val_loss: 0.3165 -
val_precision_5: 0.4516 - val_recall_5: 0.1261
Epoch 10/50
293/293 8s 27ms/step -
accuracy: 0.9160 - auc: 0.7871 - loss: 0.2559 - precision_5: 0.7525 - recall_5:
0.2025 - val_accuracy: 0.9042 - val_auc: 0.6471 - val_loss: 0.3129 -
val_precision_5: 0.4706 - val_recall_5: 0.0721
Epoch 11/50
293/293 8s 27ms/step -
accuracy: 0.9233 - auc: 0.7854 - loss: 0.2411 - precision_5: 0.8325 - recall_5:
0.1990 - val_accuracy: 0.8930 - val_auc: 0.6318 - val_loss: 0.3352 -
val_precision_5: 0.3108 - val_recall_5: 0.1036
Epoch 12/50
293/293 8s 28ms/step -
accuracy: 0.9190 - auc: 0.7971 - loss: 0.2461 - precision_5: 0.8162 - recall_5:
0.2034 - val_accuracy: 0.8990 - val_auc: 0.6532 - val_loss: 0.3202 -
val_precision_5: 0.3793 - val_recall_5: 0.0991
Epoch 13/50
293/293 8s 27ms/step -
accuracy: 0.9164 - auc: 0.8110 - loss: 0.2467 - precision_5: 0.8274 - recall_5:
0.1970 - val_accuracy: 0.8999 - val_auc: 0.6173 - val_loss: 0.3404 -
val_precision_5: 0.4091 - val_recall_5: 0.1216
Epoch 14/50
293/293 8s 27ms/step -
accuracy: 0.9217 - auc: 0.8096 - loss: 0.2387 - precision_5: 0.8085 - recall_5:
0.2482 - val_accuracy: 0.8969 - val_auc: 0.6337 - val_loss: 0.3464 -
val_precision_5: 0.3662 - val_recall_5: 0.1171
Epoch 15/50
293/293 8s 28ms/step -
accuracy: 0.9226 - auc: 0.8157 - loss: 0.2353 - precision_5: 0.7771 - recall_5:
0.2612 - val_accuracy: 0.8964 - val_auc: 0.6266 - val_loss: 0.3502 -
val_precision_5: 0.3718 - val_recall_5: 0.1306
Epoch 16/50
293/293 8s 27ms/step -
accuracy: 0.9210 - auc: 0.8437 - loss: 0.2281 - precision_5: 0.8037 - recall_5:
0.2550 - val_accuracy: 0.8990 - val_auc: 0.6098 - val_loss: 0.3654 -
val_precision_5: 0.3906 - val_recall_5: 0.1126
Epoch 17/50
293/293 8s 28ms/step -
accuracy: 0.9278 - auc: 0.8379 - loss: 0.2197 - precision_5: 0.8639 - recall_5:
0.2781 - val_accuracy: 0.8982 - val_auc: 0.6492 - val_loss: 0.3523 -
val_precision_5: 0.3974 - val_recall_5: 0.1396
Epoch 18/50

293/293 8s 28ms/step -
accuracy: 0.9269 - auc: 0.8543 - loss: 0.2167 - precision_5: 0.8728 - recall_5:
0.3078 - val_accuracy: 0.8956 - val_auc: 0.6393 - val_loss: 0.3674 -
val_precision_5: 0.3690 - val_recall_5: 0.1396
Epoch 19/50

293/293 8s 27ms/step -
accuracy: 0.9268 - auc: 0.8642 - loss: 0.2161 - precision_5: 0.7765 - recall_5:
0.2923 - val_accuracy: 0.8969 - val_auc: 0.6374 - val_loss: 0.3727 -
val_precision_5: 0.3797 - val_recall_5: 0.1351
Epoch 20/50

293/293 8s 27ms/step -
accuracy: 0.9307 - auc: 0.8611 - loss: 0.2077 - precision_5: 0.8544 - recall_5:
0.3221 - val_accuracy: 0.8935 - val_auc: 0.6497 - val_loss: 0.3901 -
val_precision_5: 0.3333 - val_recall_5: 0.1216
Epoch 21/50

293/293 8s 27ms/step -
accuracy: 0.9237 - auc: 0.8709 - loss: 0.2136 - precision_5: 0.8240 - recall_5:
0.2953 - val_accuracy: 0.8887 - val_auc: 0.6475 - val_loss: 0.3956 -
val_precision_5: 0.2683 - val_recall_5: 0.0991
Epoch 22/50

293/293 8s 27ms/step -
accuracy: 0.9284 - auc: 0.8670 - loss: 0.2087 - precision_5: 0.8691 - recall_5:
0.3124 - val_accuracy: 0.8956 - val_auc: 0.6408 - val_loss: 0.4048 -
val_precision_5: 0.3472 - val_recall_5: 0.1126
Epoch 23/50

293/293 8s 28ms/step -
accuracy: 0.9339 - auc: 0.8855 - loss: 0.1957 - precision_5: 0.9045 - recall_5:
0.3508 - val_accuracy: 0.8935 - val_auc: 0.6526 - val_loss: 0.4077 -
val_precision_5: 0.3548 - val_recall_5: 0.1486
Epoch 24/50

293/293 8s 27ms/step -
accuracy: 0.9329 - auc: 0.8883 - loss: 0.1954 - precision_5: 0.8794 - recall_5:
0.3564 - val_accuracy: 0.8990 - val_auc: 0.6386 - val_loss: 0.4112 -
val_precision_5: 0.4079 - val_recall_5: 0.1396
Epoch 25/50

293/293 8s 27ms/step -
accuracy: 0.9316 - auc: 0.8942 - loss: 0.1963 - precision_5: 0.8959 - recall_5:
0.3726 - val_accuracy: 0.8943 - val_auc: 0.6529 - val_loss: 0.4689 -
val_precision_5: 0.3656 - val_recall_5: 0.1532
Epoch 26/50

293/293 8s 27ms/step -
accuracy: 0.9344 - auc: 0.8956 - loss: 0.1918 - precision_5: 0.8574 - recall_5:
0.3741 - val_accuracy: 0.8960 - val_auc: 0.6421 - val_loss: 0.4192 -
val_precision_5: 0.3735 - val_recall_5: 0.1396
Epoch 27/50

293/293 8s 27ms/step -
accuracy: 0.9344 - auc: 0.8929 - loss: 0.1864 - precision_5: 0.8711 - recall_5:
0.3643 - val_accuracy: 0.8866 - val_auc: 0.6482 - val_loss: 0.4417 -

val_precision_5: 0.3193 - val_recall_5: 0.1712
Epoch 28/50
293/293 8s 27ms/step -
accuracy: 0.9390 - auc: 0.9118 - loss: 0.1737 - precision_5: 0.8609 - recall_5:
0.4000 - val_accuracy: 0.8879 - val_auc: 0.6633 - val_loss: 0.4549 -
val_precision_5: 0.3246 - val_recall_5: 0.1667
Epoch 29/50
293/293 8s 27ms/step -
accuracy: 0.9335 - auc: 0.9017 - loss: 0.1868 - precision_5: 0.8444 - recall_5:
0.3910 - val_accuracy: 0.8870 - val_auc: 0.6402 - val_loss: 0.4376 -
val_precision_5: 0.2941 - val_recall_5: 0.1351
Epoch 30/50
293/293 8s 28ms/step -
accuracy: 0.9349 - auc: 0.9076 - loss: 0.1835 - precision_5: 0.8390 - recall_5:
0.4348 - val_accuracy: 0.8947 - val_auc: 0.6454 - val_loss: 0.4431 -
val_precision_5: 0.3776 - val_recall_5: 0.1667
Epoch 31/50
293/293 8s 28ms/step -
accuracy: 0.9398 - auc: 0.9146 - loss: 0.1711 - precision_5: 0.8789 - recall_5:
0.4168 - val_accuracy: 0.8956 - val_auc: 0.6253 - val_loss: 0.4744 -
val_precision_5: 0.3830 - val_recall_5: 0.1622
Epoch 32/50
293/293 8s 27ms/step -
accuracy: 0.9444 - auc: 0.9191 - loss: 0.1634 - precision_5: 0.8868 - recall_5:
0.4435 - val_accuracy: 0.8964 - val_auc: 0.6307 - val_loss: 0.4957 -
val_precision_5: 0.3936 - val_recall_5: 0.1667
Epoch 33/50
293/293 8s 27ms/step -
accuracy: 0.9397 - auc: 0.9178 - loss: 0.1690 - precision_5: 0.8659 - recall_5:
0.4221 - val_accuracy: 0.8900 - val_auc: 0.6382 - val_loss: 0.5027 -
val_precision_5: 0.3451 - val_recall_5: 0.1757
Epoch 34/50
293/293 8s 28ms/step -
accuracy: 0.9405 - auc: 0.9257 - loss: 0.1677 - precision_5: 0.8754 - recall_5:
0.4732 - val_accuracy: 0.8840 - val_auc: 0.6366 - val_loss: 0.5038 -
val_precision_5: 0.3287 - val_recall_5: 0.2117
Epoch 35/50
293/293 8s 27ms/step -
accuracy: 0.9418 - auc: 0.9207 - loss: 0.1669 - precision_5: 0.8682 - recall_5:
0.4875 - val_accuracy: 0.8952 - val_auc: 0.6200 - val_loss: 0.5509 -
val_precision_5: 0.3789 - val_recall_5: 0.1622
Epoch 36/50
293/293 8s 27ms/step -
accuracy: 0.9415 - auc: 0.9112 - loss: 0.1693 - precision_5: 0.8428 - recall_5:
0.4626 - val_accuracy: 0.8875 - val_auc: 0.6469 - val_loss: 0.5060 -
val_precision_5: 0.3481 - val_recall_5: 0.2117
Epoch 37/50
293/293 8s 27ms/step -

accuracy: 0.9443 - auc: 0.9203 - loss: 0.1588 - precision_5: 0.8744 - recall_5: 0.4524 - val_accuracy: 0.8973 - val_auc: 0.6386 - val_loss: 0.4936 - val_precision_5: 0.4135 - val_recall_5: 0.1937
Epoch 38/50
293/293 8s 28ms/step -
accuracy: 0.9454 - auc: 0.9268 - loss: 0.1581 - precision_5: 0.8818 - recall_5: 0.4686 - val_accuracy: 0.8956 - val_auc: 0.6482 - val_loss: 0.4939 - val_precision_5: 0.4018 - val_recall_5: 0.2027
Epoch 39/50
293/293 8s 28ms/step -
accuracy: 0.9485 - auc: 0.9378 - loss: 0.1476 - precision_5: 0.8863 - recall_5: 0.5135 - val_accuracy: 0.8887 - val_auc: 0.6444 - val_loss: 0.5439 - val_precision_5: 0.3443 - val_recall_5: 0.1892
Epoch 40/50
293/293 8s 27ms/step -
accuracy: 0.9438 - auc: 0.9380 - loss: 0.1523 - precision_5: 0.8730 - recall_5: 0.4880 - val_accuracy: 0.8930 - val_auc: 0.6540 - val_loss: 0.4868 - val_precision_5: 0.3704 - val_recall_5: 0.1802
Epoch 41/50
293/293 8s 27ms/step -
accuracy: 0.9464 - auc: 0.9331 - loss: 0.1549 - precision_5: 0.9044 - recall_5: 0.4901 - val_accuracy: 0.8956 - val_auc: 0.6459 - val_loss: 0.4923 - val_precision_5: 0.3854 - val_recall_5: 0.1667
Epoch 42/50
293/293 8s 28ms/step -
accuracy: 0.9530 - auc: 0.9474 - loss: 0.1372 - precision_5: 0.9023 - recall_5: 0.5646 - val_accuracy: 0.8905 - val_auc: 0.6279 - val_loss: 0.5508 - val_precision_5: 0.3482 - val_recall_5: 0.1757
Epoch 43/50
293/293 8s 27ms/step -
accuracy: 0.9438 - auc: 0.9455 - loss: 0.1483 - precision_5: 0.8872 - recall_5: 0.4983 - val_accuracy: 0.8879 - val_auc: 0.6289 - val_loss: 0.5745 - val_precision_5: 0.3182 - val_recall_5: 0.1577
Epoch 44/50
293/293 8s 27ms/step -
accuracy: 0.9518 - auc: 0.9428 - loss: 0.1411 - precision_5: 0.9323 - recall_5: 0.5328 - val_accuracy: 0.8892 - val_auc: 0.6224 - val_loss: 0.5875 - val_precision_5: 0.3333 - val_recall_5: 0.1667
Epoch 45/50
293/293 8s 28ms/step -
accuracy: 0.9462 - auc: 0.9412 - loss: 0.1489 - precision_5: 0.8961 - recall_5: 0.5054 - val_accuracy: 0.8939 - val_auc: 0.6340 - val_loss: 0.5849 - val_precision_5: 0.3587 - val_recall_5: 0.1486
Epoch 46/50
293/293 8s 28ms/step -
accuracy: 0.9533 - auc: 0.9454 - loss: 0.1358 - precision_5: 0.9262 - recall_5: 0.5561 - val_accuracy: 0.8883 - val_auc: 0.6283 - val_loss: 0.5643 - val_precision_5: 0.3415 - val_recall_5: 0.1892


```

Epoch 47/50
293/293      8s 28ms/step -
accuracy: 0.9532 - auc: 0.9465 - loss: 0.1347 - precision_5: 0.8708 - recall_5:
0.5730 - val_accuracy: 0.8870 - val_auc: 0.6269 - val_loss: 0.5829 -
val_precision_5: 0.3279 - val_recall_5: 0.1802
Epoch 48/50
293/293      8s 28ms/step -
accuracy: 0.9516 - auc: 0.9491 - loss: 0.1386 - precision_5: 0.9053 - recall_5:
0.5477 - val_accuracy: 0.8806 - val_auc: 0.6411 - val_loss: 0.6040 -
val_precision_5: 0.3161 - val_recall_5: 0.2207
Epoch 49/50
293/293      8s 27ms/step -
accuracy: 0.9483 - auc: 0.9511 - loss: 0.1427 - precision_5: 0.8736 - recall_5:
0.5566 - val_accuracy: 0.8947 - val_auc: 0.6112 - val_loss: 0.6340 -
val_precision_5: 0.3846 - val_recall_5: 0.1802
Epoch 50/50
293/293      8s 27ms/step -
accuracy: 0.9472 - auc: 0.9459 - loss: 0.1450 - precision_5: 0.8952 - recall_5:
0.5290 - val_accuracy: 0.8922 - val_auc: 0.6244 - val_loss: 0.5922 -
val_precision_5: 0.3636 - val_recall_5: 0.1802

```

```

[331]: num_clients, num_months, num_features = X_test.shape

# Reshape to 2D for scaling
X_2d = X_test.reshape(-1, num_features) # shape: (num_clients * num_months,
    ↪ num_features)

# Fit scaler on training data only!
X_2d_scaled = scaler.transform(X_2d)

# Reshape back to 3D
X_test_scaled = X_2d_scaled.reshape(num_clients, num_months, num_features)

```

```

[335]: Y_test.mean()

```

```

[335]: np.float64(0.10130047912388775)

```

```

[337]: y_pred = (model.predict(X_test_scaled) > 0.5).astype(int)
y_pred_proba = model.predict(X_test_scaled)

fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
roc_auc = roc_auc_score(Y_test, y_pred_proba)


# Plot ROC curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')

```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for XGBoost Churn Model')
plt.legend(loc='lower right')
plt.show()

print(f"ROC AUC Score: {roc_auc:.2f}")
```

92/92 1s 7ms/step
 92/92 1s 7ms/step



ROC AUC Score: 0.61

```
[338]: # Classification report
print(classification_report(Y_test, y_pred))

# Confusion matrix
cm = confusion_matrix(Y_test, y_pred)
print("Confusion Matrix:\n", cm)
print(f'F1 Score: {f1_score(Y_test, y_pred)}')
```

	precision	recall	f1-score	support
0.0	0.92	0.97	0.94	2626
1.0	0.45	0.20	0.28	296
accuracy			0.89	2922
macro avg	0.68	0.59	0.61	2922
weighted avg	0.87	0.89	0.88	2922

Confusion Matrix:

```
[[2554  72]
```

```
[ 237  59]]
```

F1 Score: 0.27634660421545665

1.5 Conclusion

BiDirectional LSTM-CNN Model worked better than DNN we covered before. Summary for metrics for three models:

Model Performance Comparison

Model	F1 Score	Recall	ROC-AUC	Accuracy
BiLSTM-CNN	0.28	0.20	0.61	0.89
DNN	0.25	0.19	0.62	0.89
XGBoost	0.27	0.50	0.69	0.74

BiLSTM-CNN has highest F1 score among the three models we tested, however it has low AUC score and recall compared to Gradient Boost model. Generally, Gradient Boosting is considered industry standard for churn analysis and through this result we can observe that it is indeed better in getting higher recall which is essential for churn prediction.

This project was started with the client hypothesis that Price is a major factor leading to churn. Since XGBoost Model is best model here we used feature importance for that model and found that sale promotion id or group is most important factor leading to churn, solidifying our findings during EDA that proportion of churn for some subgroups of clients were higher than usual. Price was a reasonable factor, but not a major one as even margin is better indicator.

We got best model F1 score at 0.27 and recall of 0.50 with ROC-AUC score at 0.69. While accuracy is not important its 0.74 for XGBoost model, its not high but since we focused on improving ROC scores on all models it was acceptable trade. All of our models worked a lot better than BCGX model answer where we had F1 Score of 0.09, recall of 0.05, and ROC score of 0.52.

We can get better model performance but for that we need more data and better feature engineering which is not possible with the dataset in hand. Also, while both DNN Model and BiLSTM-CNN model underperformed compared to Gradient Boost model because they do not work good with such tabular data in general, it is to be noted that in scenario where we have a large amount of

data both of these models could work better. And we can even use a mix of Gradient Boosting and BiLSTM-CNN models by using ensemble methods on BiLSTM-CNN model.